

Getting Started with MLflow

Understanding MLflow



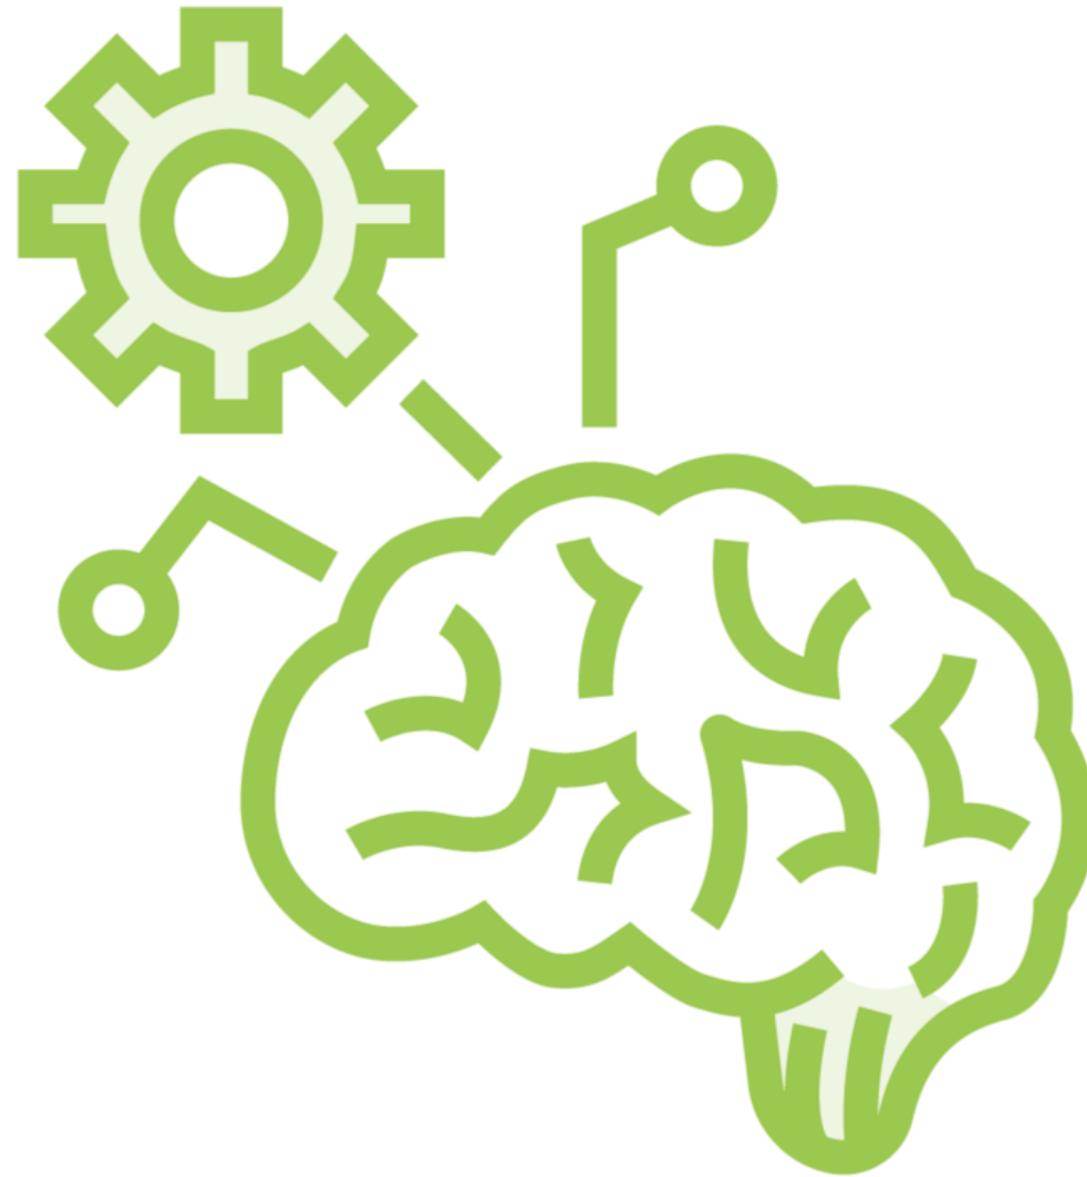
Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Why MLflow?



Machine learning is everywhere

Growth of ML has created new challenges

Need for new tools gave rise to MLOps

MLflow solves certain challenges that MLOps aims to address



Basics



Instructor (me)

Target audience (you)

Prerequisites

Demo story description

MLflow and problems it solves



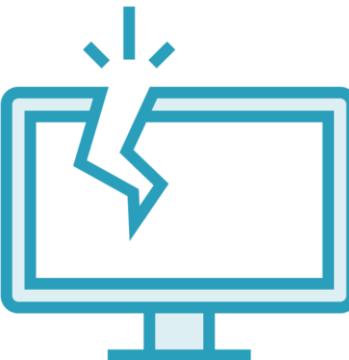
Me



Passionate about data and machine learning



Fraud detection



Experienced first-hand the challenges I talk about



You



Data scientist
Machine learning engineer
Hobbyist
Engineering manager
Infrastructure engineer
Anything related to ML



Prerequisites

Comfortable with Python
**Minimal practical experience with building
ML models**



Story

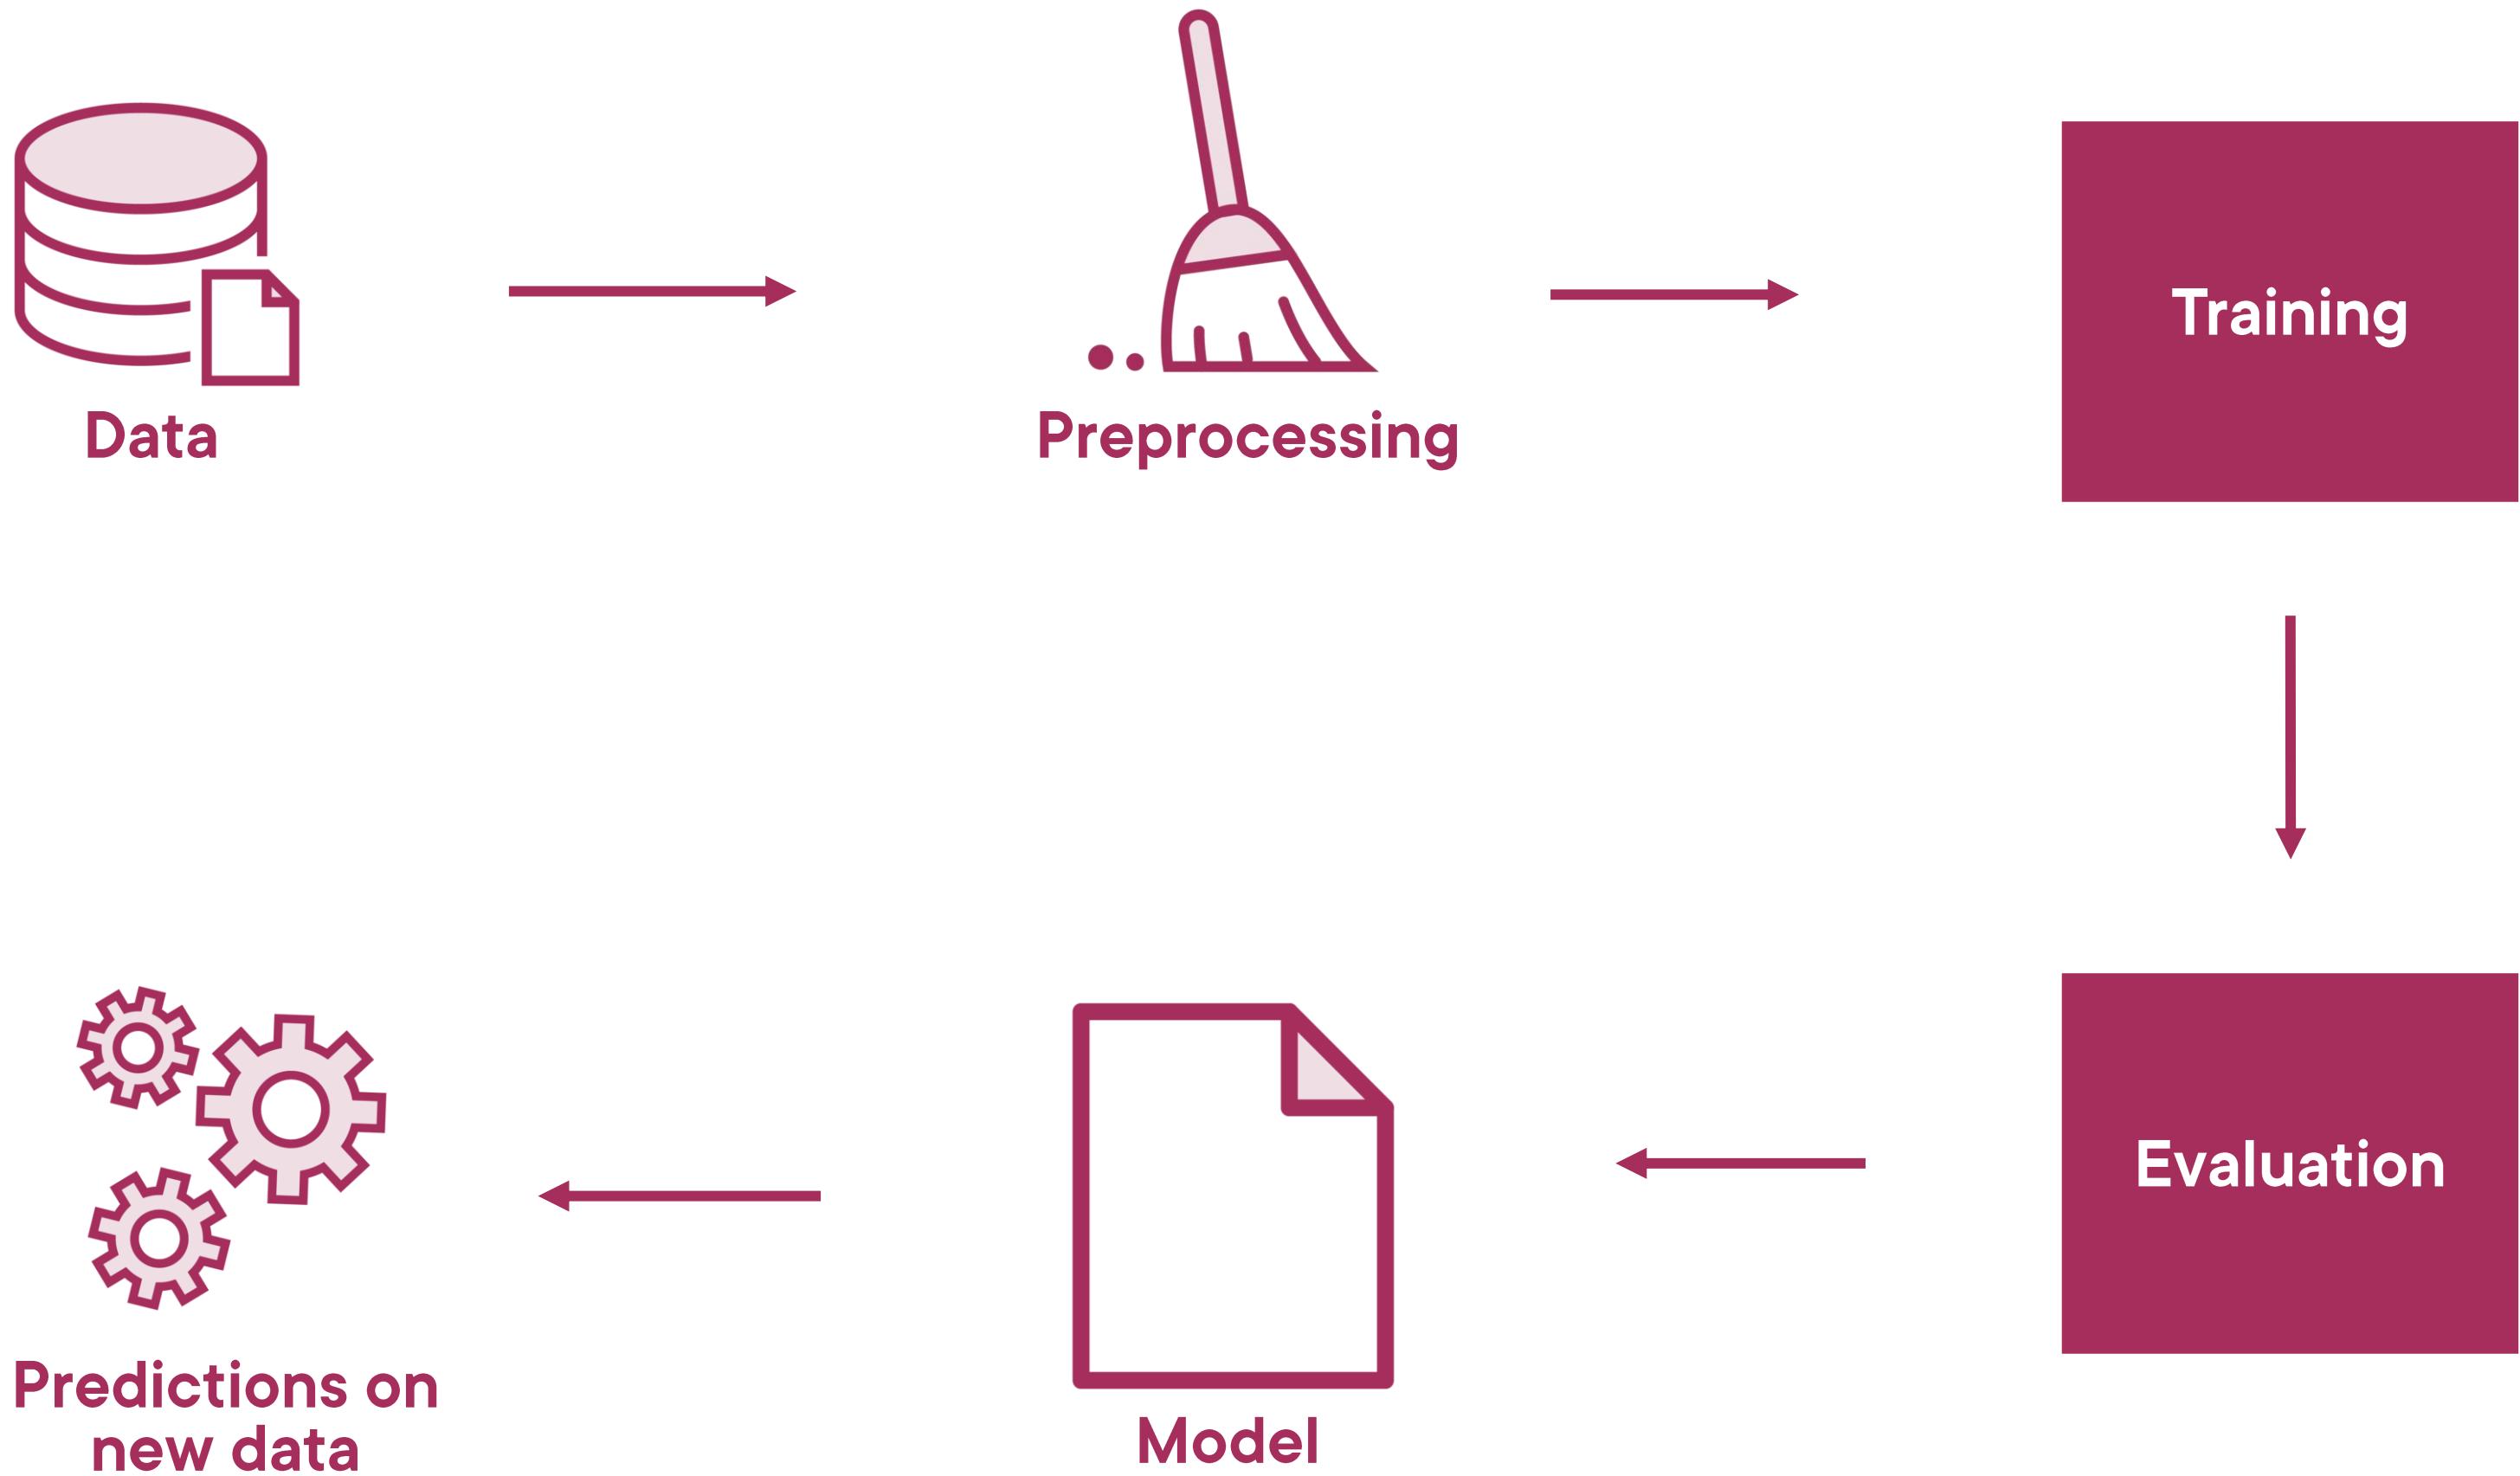


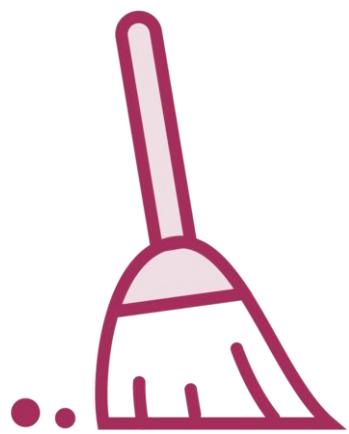
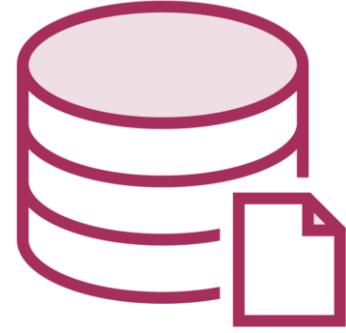
Real estate company
One of the first data scientists
Prediction of house prices



What Is MLflow?

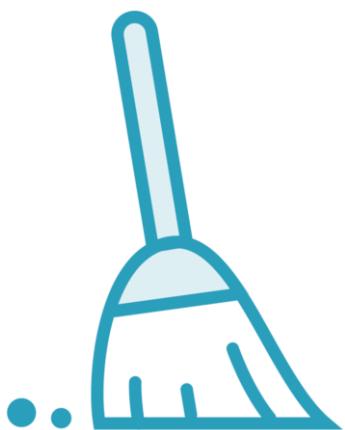




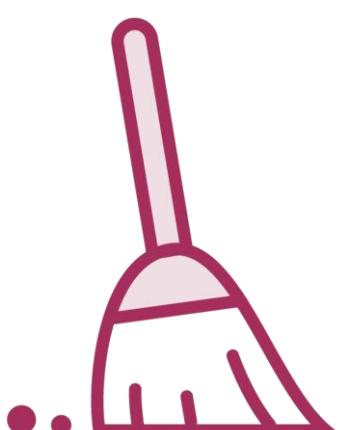


Training

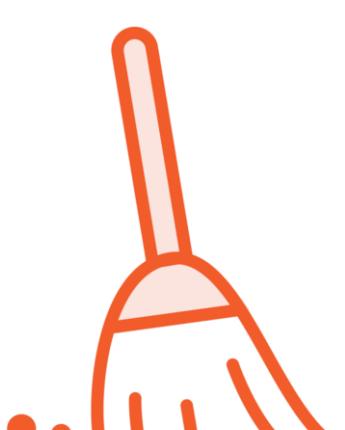




Training



Training

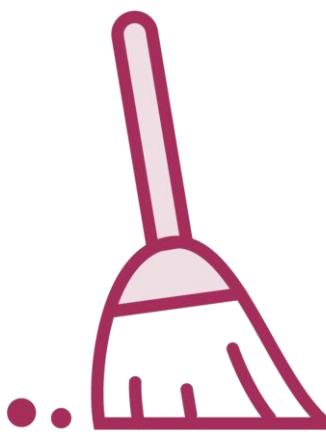
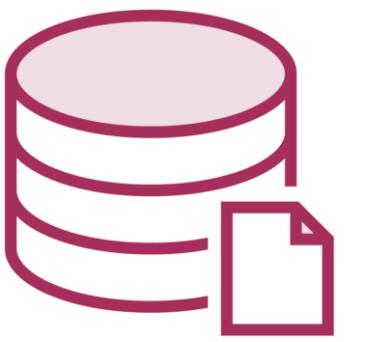


Training

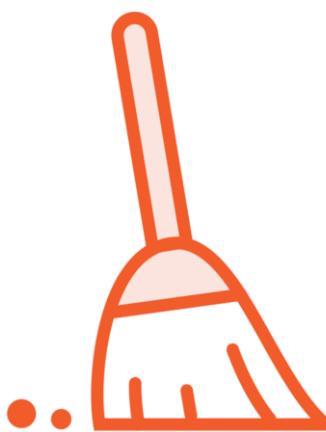




Training



Training



Training



Tracking the Changes



Reproducibility



Traceability



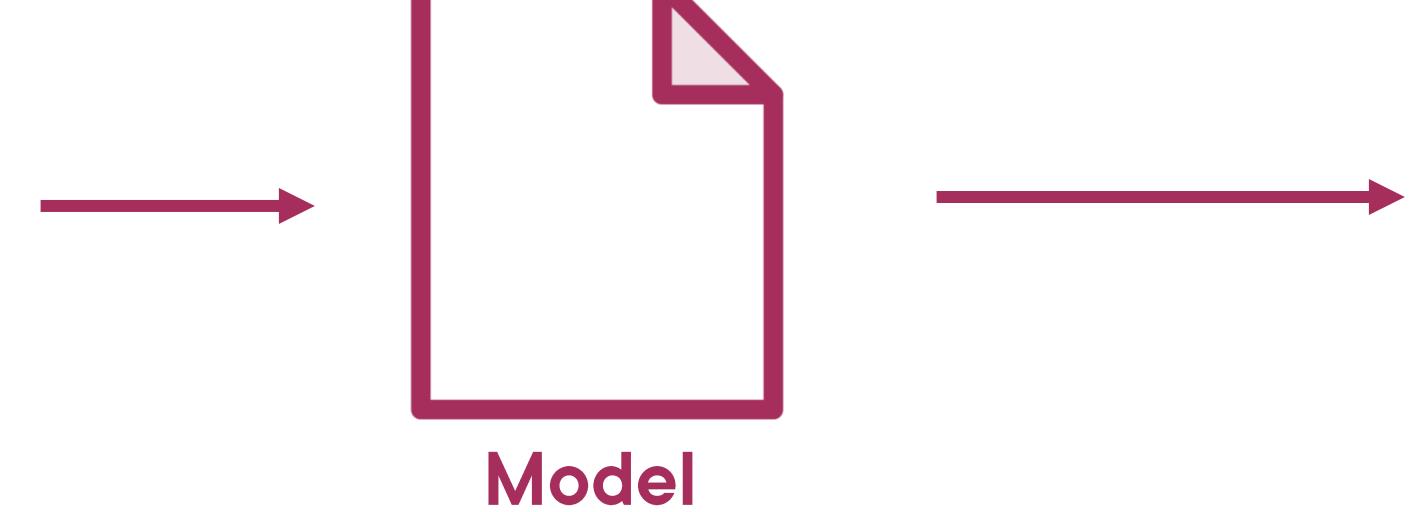
Notes and spreadsheets are not scalable

MLflow Tracking

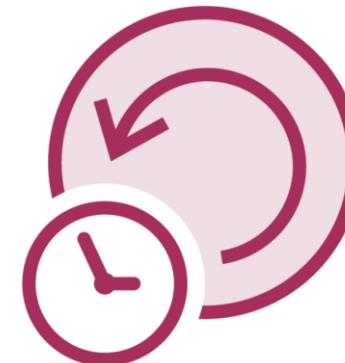
- Stores data for all experiments
- User interface
- Programmatic querying



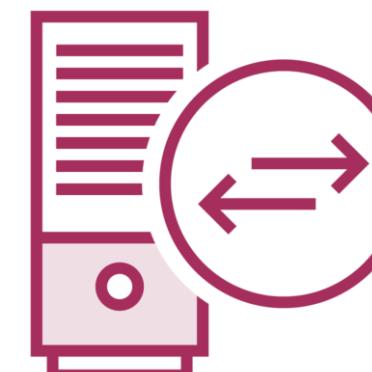
Training and evaluation



Production environment



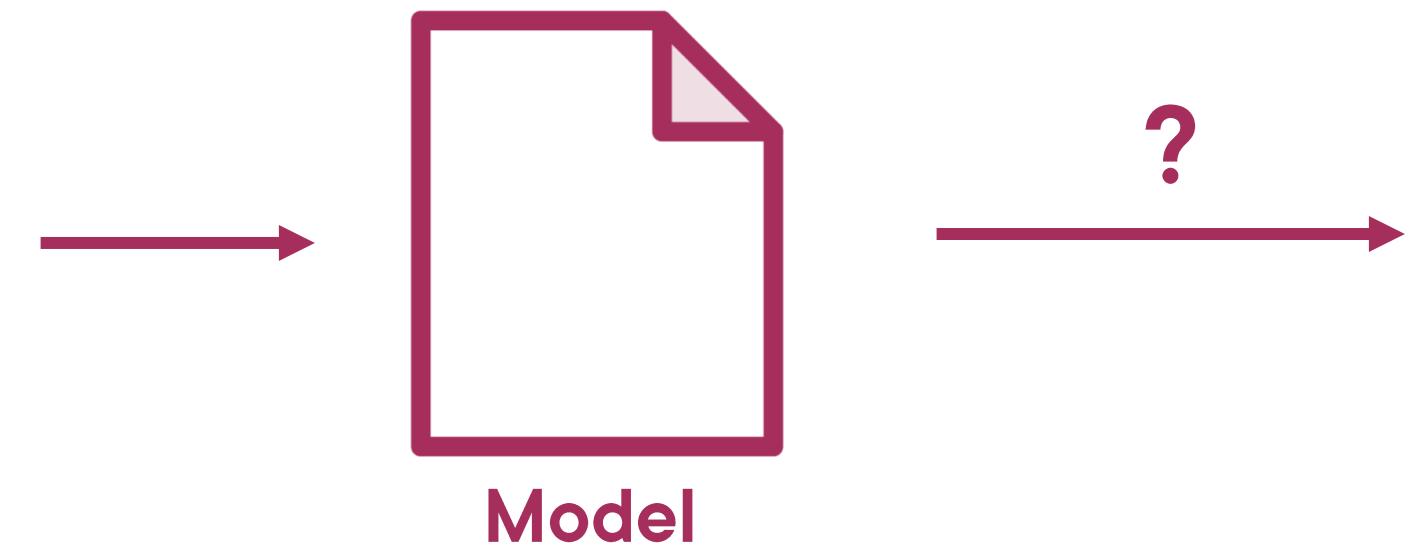
Batch predictions



Real-time predictions



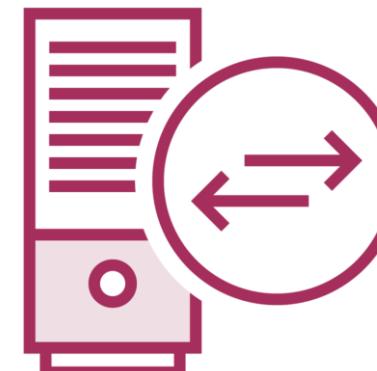
Training and evaluation



Production environment



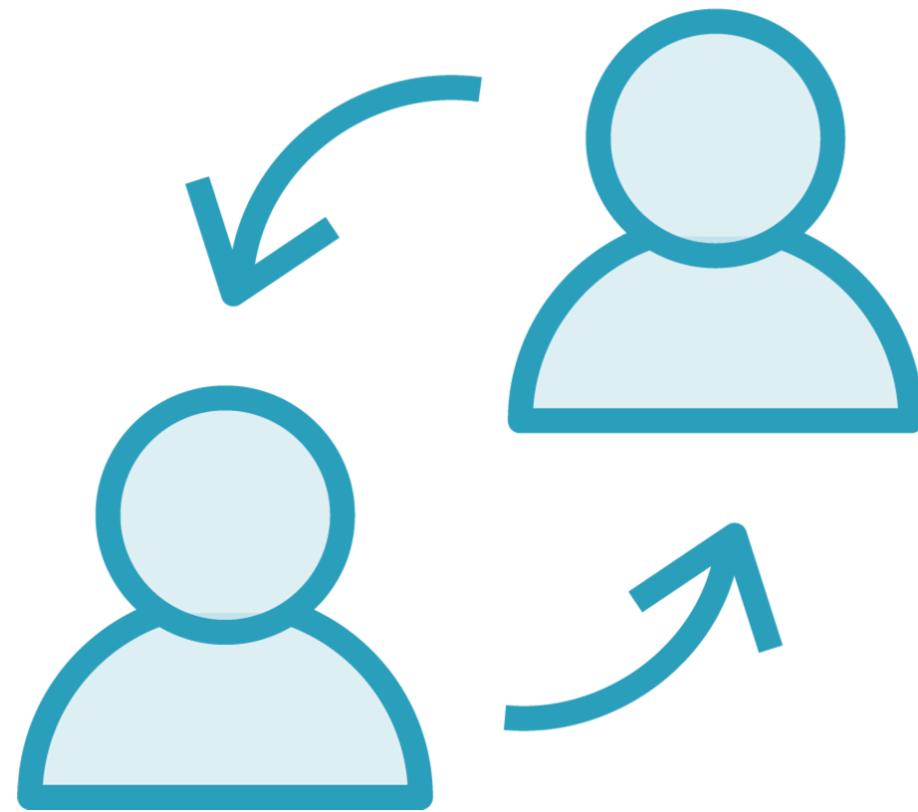
Batch predictions



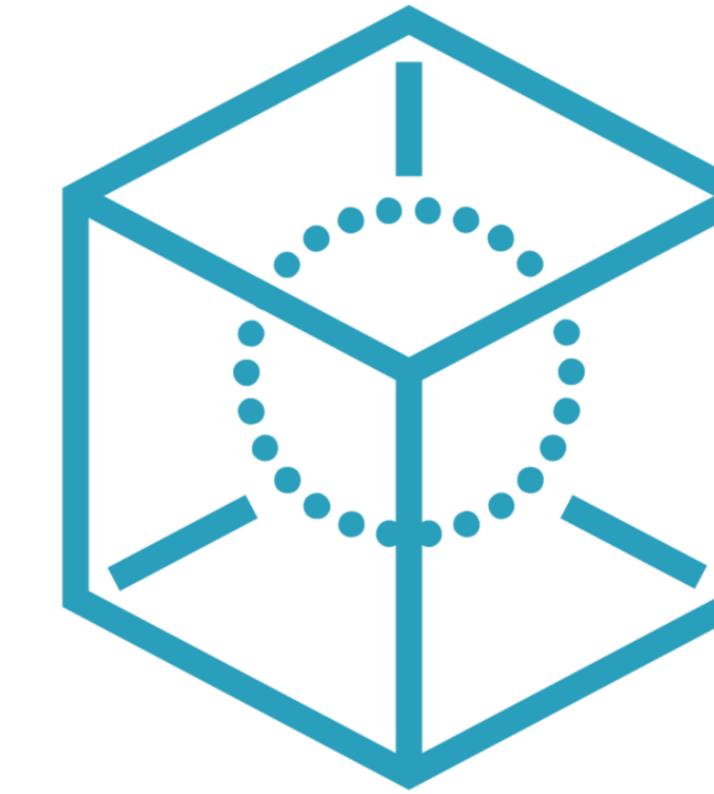
Real-time predictions



Model Handover



Extensive explanations required



Additional layers of code



MLflow Models

Format for packing models

Running predictions against existing dataset

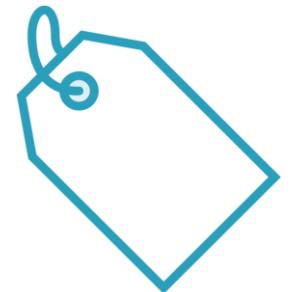
Real-time serving



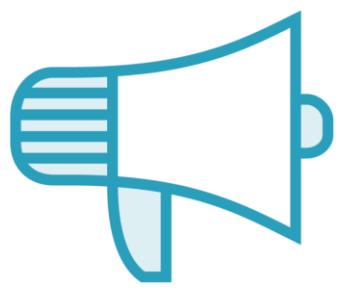
Managing Models



Model discovery



Model metadata management



Traditionally requires explicit communication



Human error likely to happen



MLflow Model Registry

Centralized model catalogue

**Models can be used/deployed directly from
the registry**

**Can be used as a base for building automated
tools**



MLflow Projects

Another MLflow component

**Opinionated – often not compatible with
existing choice of tools inside ML teams**

Won't be explored in this course



MLflow

Tracking

Models

Model Registry



Summary



Ubiquity of machine learning

Target audience and prerequisites

MLflow components



Tracking ML Experiments



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Overview



Dataset

Model training pipeline

Vocabulary

MLflow Tracking: Demo



Dataset



Task

Predict market price of a property given some of its characteristics



Ames Housing Dataset



De Cock D. 2011. Ames, Iowa: Alternative to the Boston Housing Data as an End of Semester Regression Project. Journal of Statistics Education; 19(3).



Complete dataset documentation:
<http://jse.amstat.org/v19n3/decock/DataDocumentation.txt>



Ames Housing Dataset

2930 rows and 83 columns

A lot of features!

Selected features:

- 'Lot Area' – size of the lot in feet
- 'Gr Liv Area' – living area above the ground level
- 'Garage Area' – garage size in square feet
- 'Bldg Type' – building type, e.g. '1Fam'

Target: 'SalePrice' – in USD



Lot Area	Gr Living Area	Garage Area	Bldg Type	SalePrice
31770	1656	528.0	1Fam	215000
11622	896	730.0	1Fam	105000
14267	1329	312.0	1Fam	172000



One-hot Encoding

	BldgType
0	1Fam
1	Duplx
2	Twnhsl
3	Duplx



	1Fam	Duplx	Twnhsl
0	1	0	0
1	0	1	0
2	0	0	1
3	0	1	0



Experimenting with ML Models



Preprocessing



Training



Preprocessing
(One-hot encoding)



Training
(Linear regression)



Evaluation
(Mean squared error)



Drop column



**Preprocessing
(One-hot encoding)**



**Training
(Linear regression)**



**Evaluation
(Mean squared error)**



Drop column:
'Lot Area'

Preprocessing
(One-hot encoding)

Training
(Linear regression)

Evaluation
(Mean squared error)

Drop column:
'Garage Area'

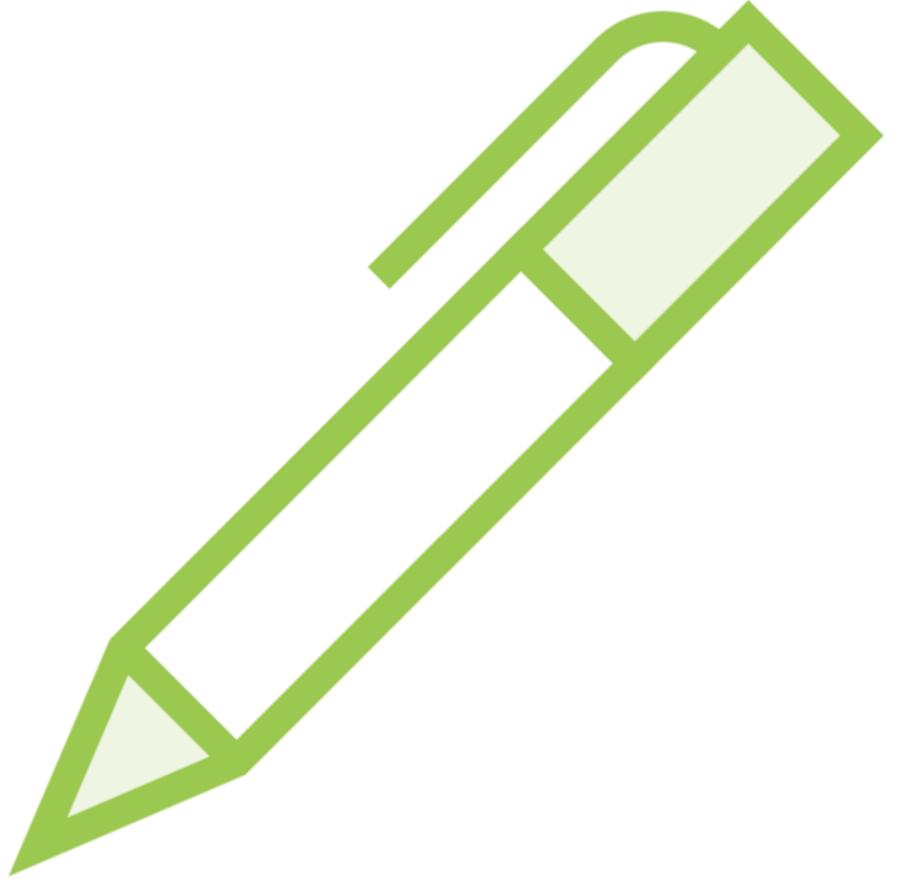
Preprocessing
(One-hot encoding)

Training
(Linear regression)

Evaluation
(Mean squared error)

Repeat for all feature columns





Initial training – error: 123.65

Dropped column: Lot Area – error: 154.34

Dropped column: Garage Area – error: 160.96



Name	Column to drop	Result
Initial training	-	123.654
Dropped column	Lot area	154.342
Dropped column	Garage area	160.964



Name	Result
Initial training	123.654

Name	Column to drop	Result
Dropped column	Lot Area	154.342
Dropped column	Garage Area	160.964



Experiment

Name: Dropped column

Column to drop	Result
Lot Area	154.342
Garage Area	160.964

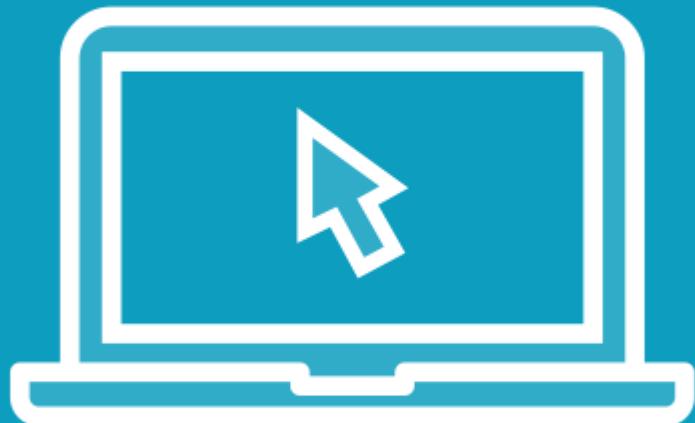
← Run

Parameter

Metric



Demo



Simple ML model for house price predictions

Creating MLflow experiments and runs

Logging parameters and metrics

Using the MLflow UI



Summary



Dataset

Training process

Concepts

- Experiment
- Run
- Parameter
- Metric

MLflow Python library

MLflow UI



Exporting Artifacts



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Overview



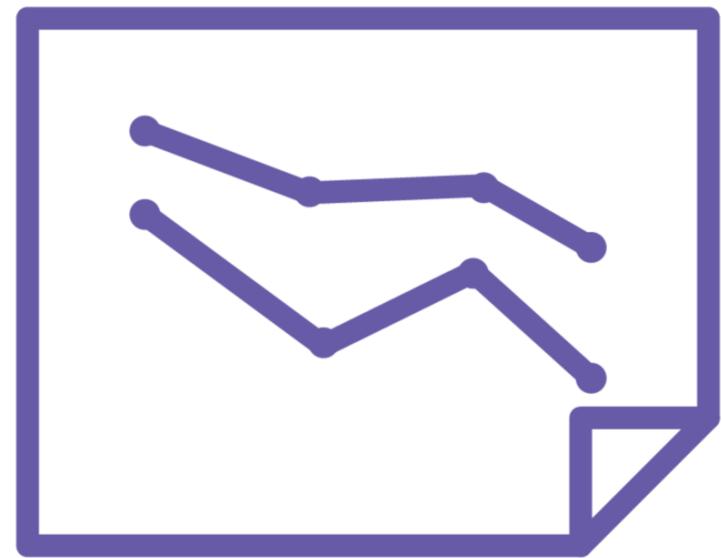
Examples of artifacts

How to save an artifact

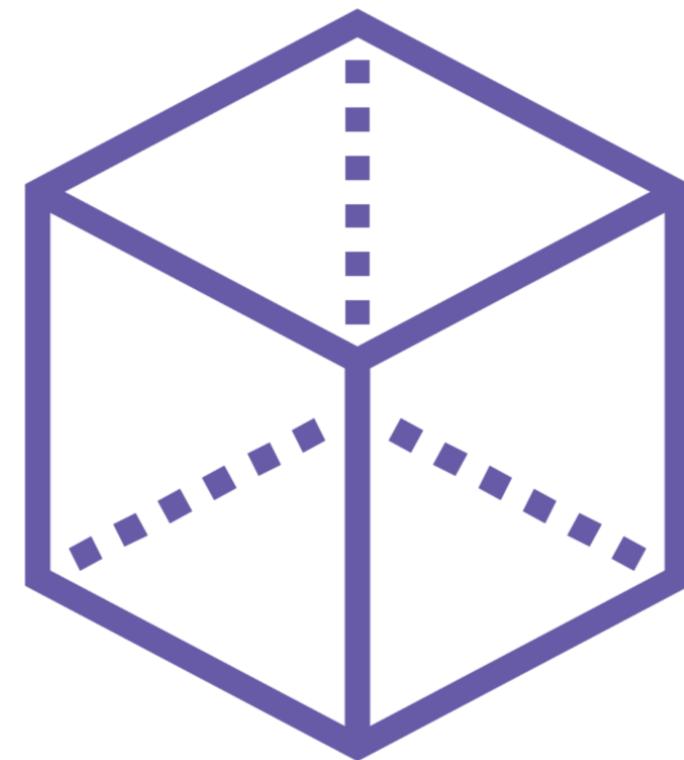
Demo



Some Examples



Plots



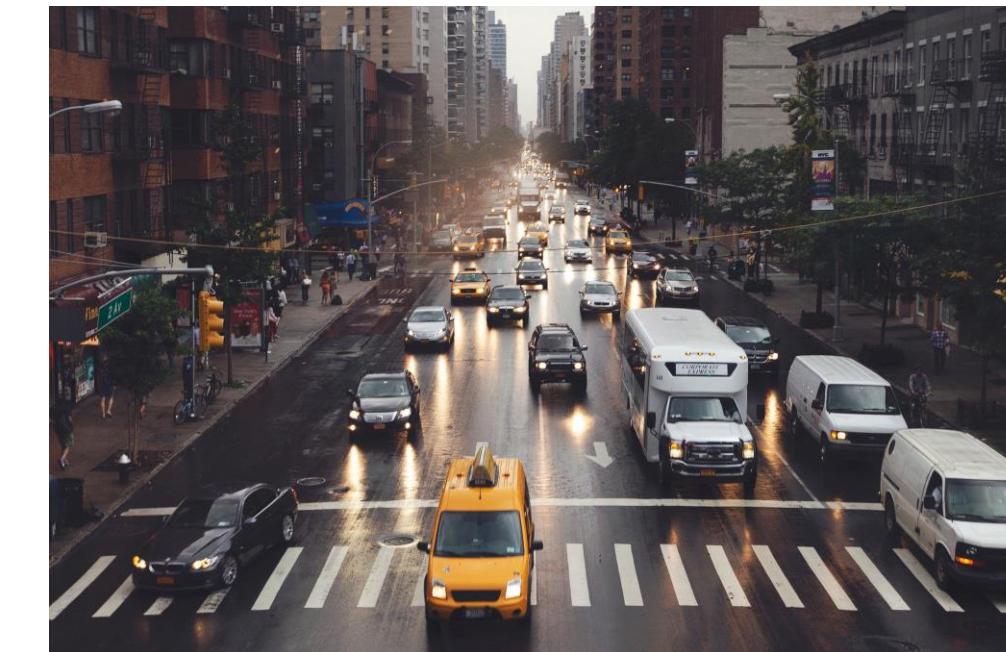
Models



Datasets



Verifying Correctness of a Transformation

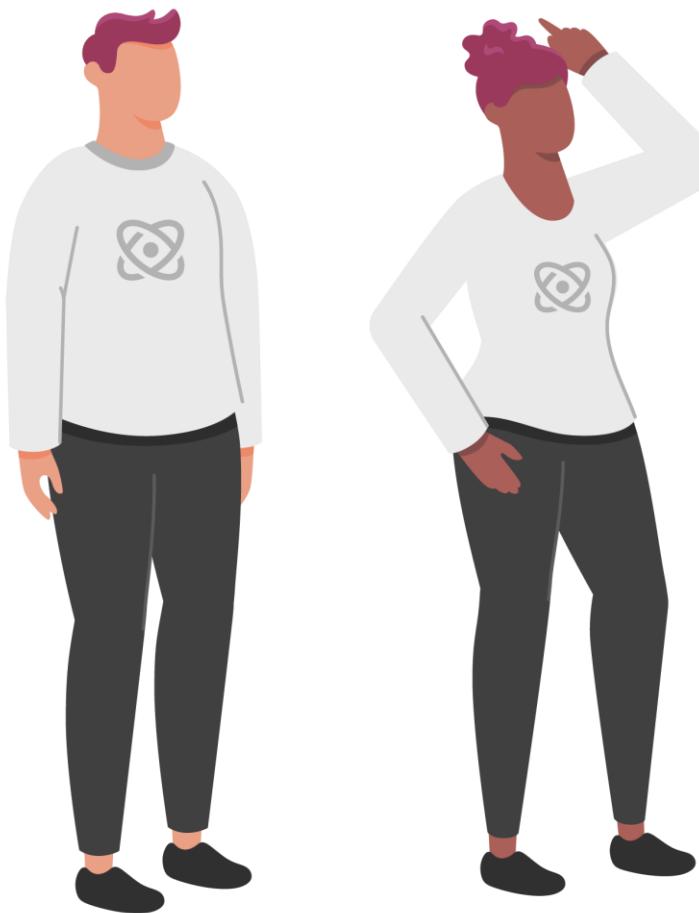




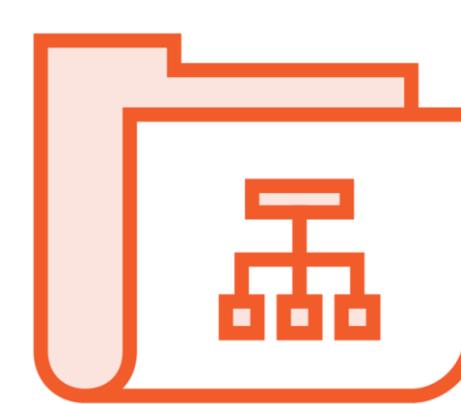
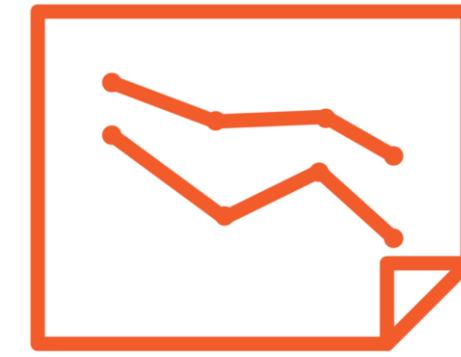
Aside from parameters and metrics, we need to store files.



Without MLflow



Everyone needs to agree on and adhere to the convention.



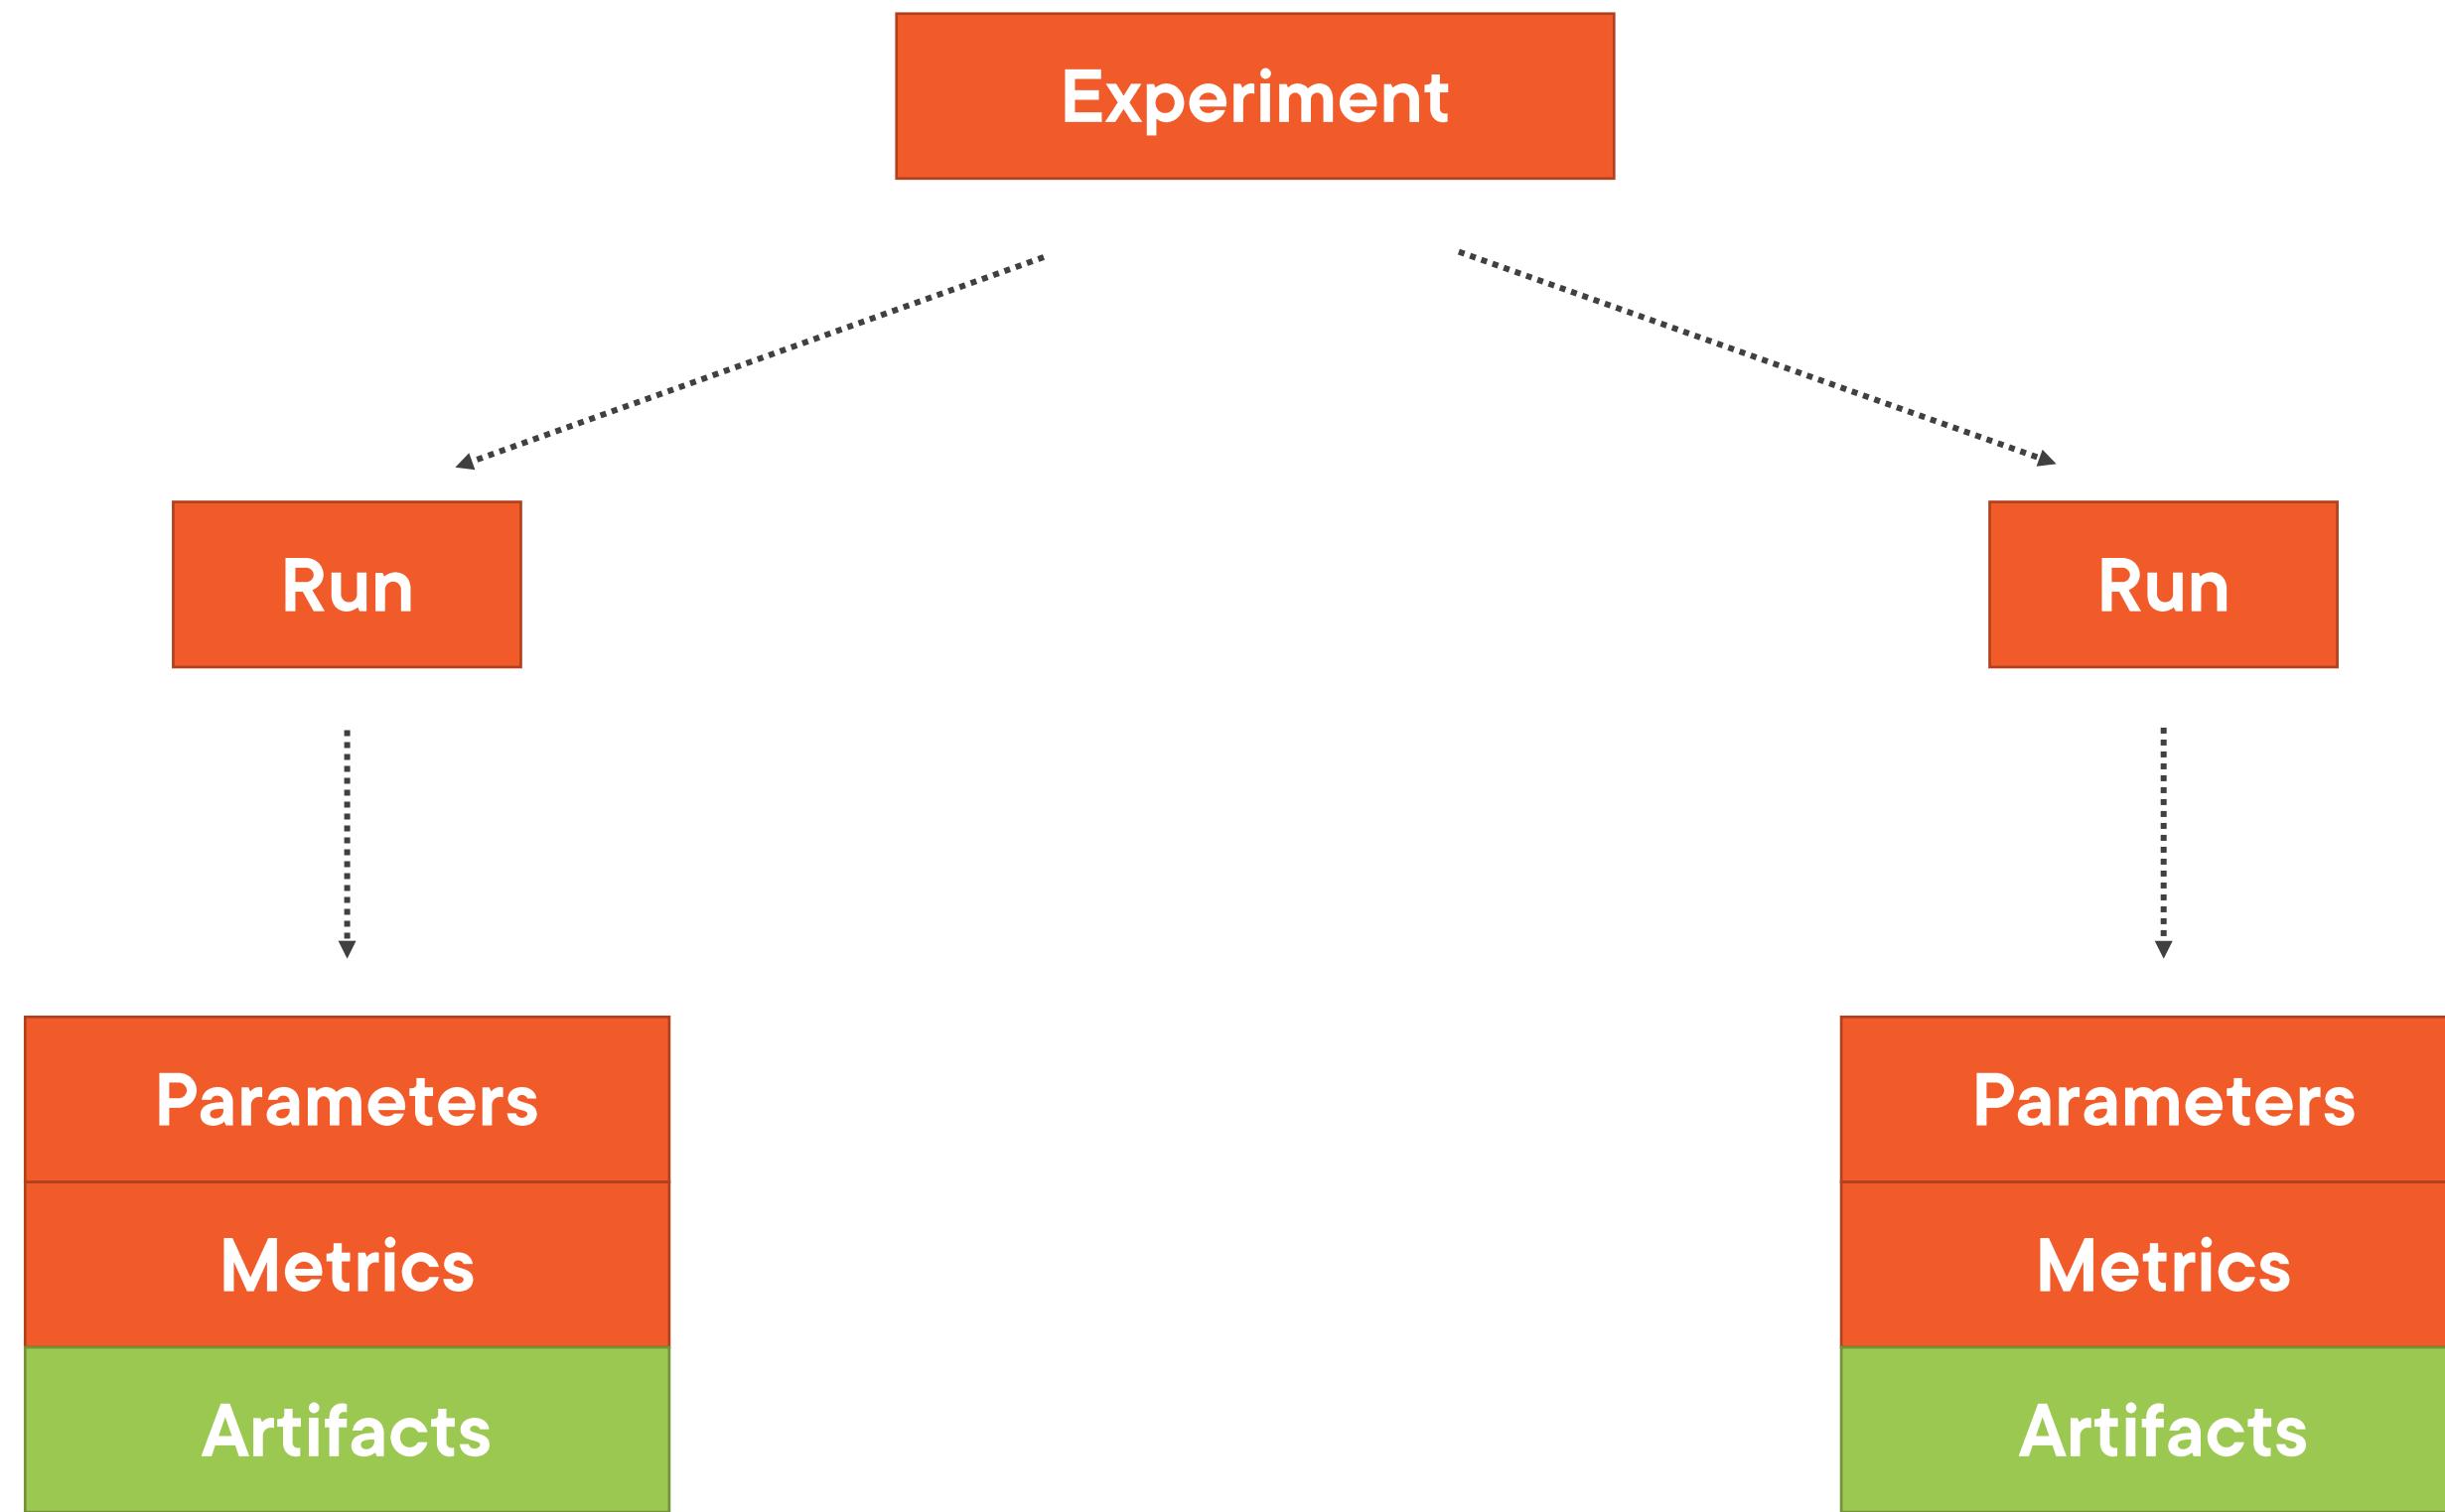
Need to manually maintain a directory structure that can be matched to experiments and runs.



With MLflow



With MLflow



Exporting an Artifact

Store the file in a temporary location on disk
Log (export) it using MLflow API



Exporting an Artifact

```
import mlflow
import pickle

model = ... # e.g. LinearRegression
serialized_model = pickle.dumps(model)

with open("some_temporary_folder/model.pkl", "wb") as m:
    m.write(serialized_model)

mlflow.log_artifact("some_temporary_folder/model.pkl")
```



Exporting an Artifact

```
import mlflow
import pickle

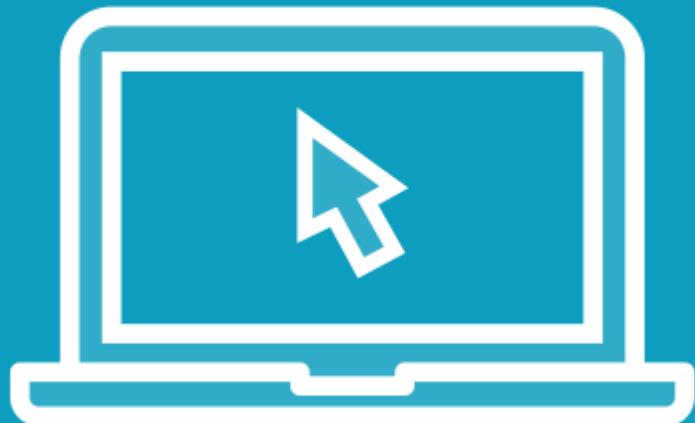
model = ... # e.g. LinearRegression
serialized_model = pickle.dumps(model)

with open("some_temporary_folder/model.pkl", "wb") as m:
    m.write(serialized_model)

mlflow.log_artifacts("some_temporary_folder")
```



Demo



Logging the model

Adding more artifacts

Browsing files using the UI



Summary



The need for artifacts in MLflow

Logging (exporting) artifacts with MLflow

Using the UI to browse the artifacts



Using MLflow in a Collaborative Scenario



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Overview



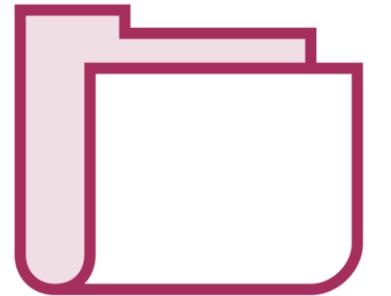
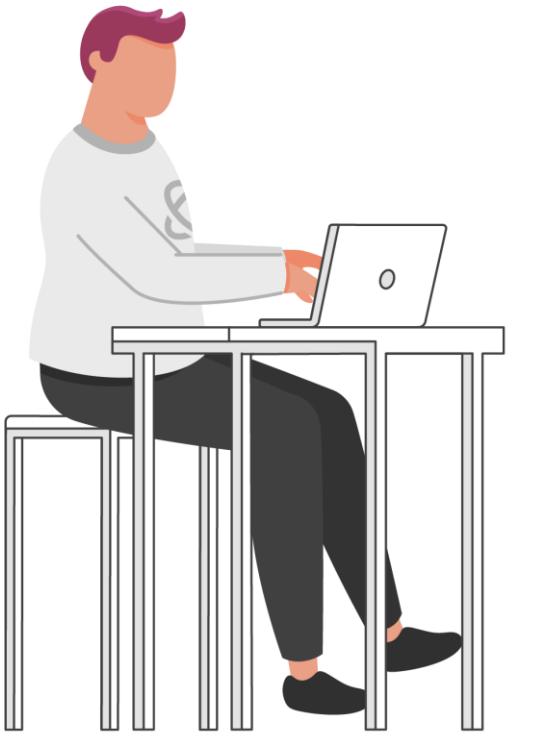
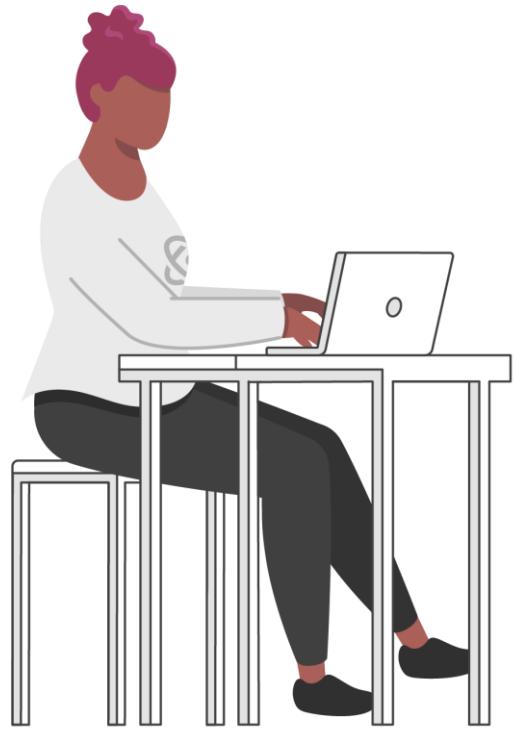
The need for a collaborative ML environment

Closer look at the current setup

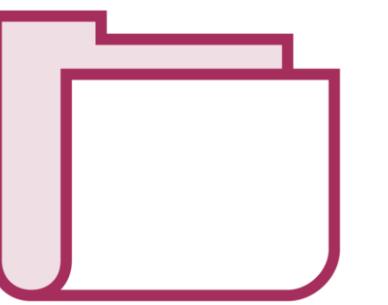
Individual components

Complete MLflow deployment



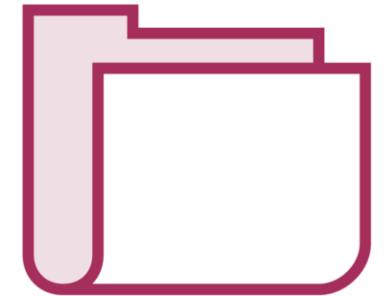
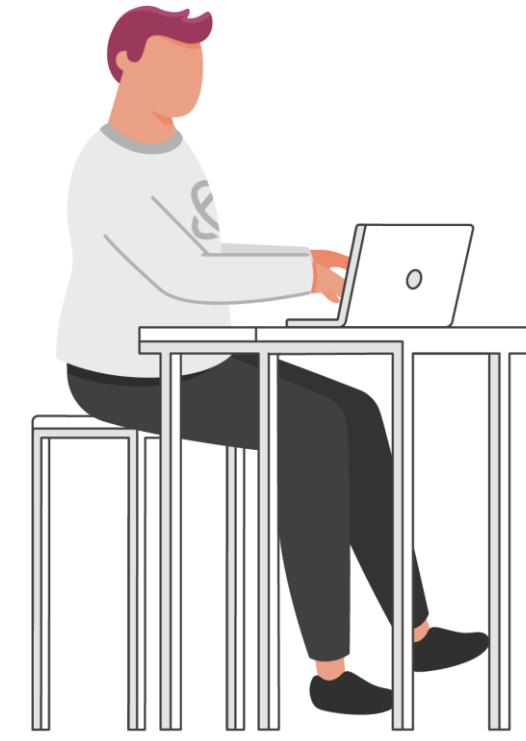
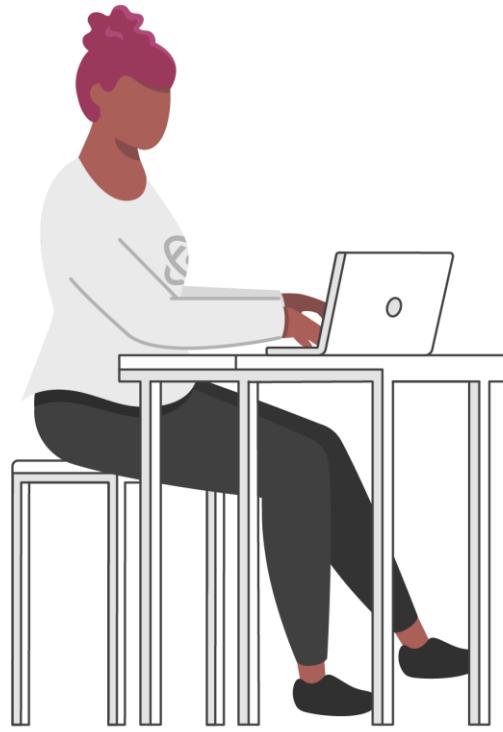


.../mlruns

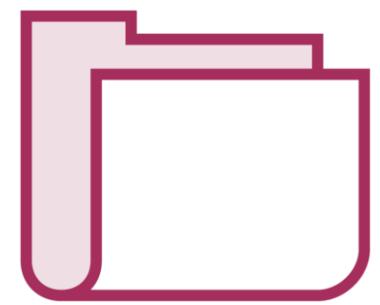


.../mlruns

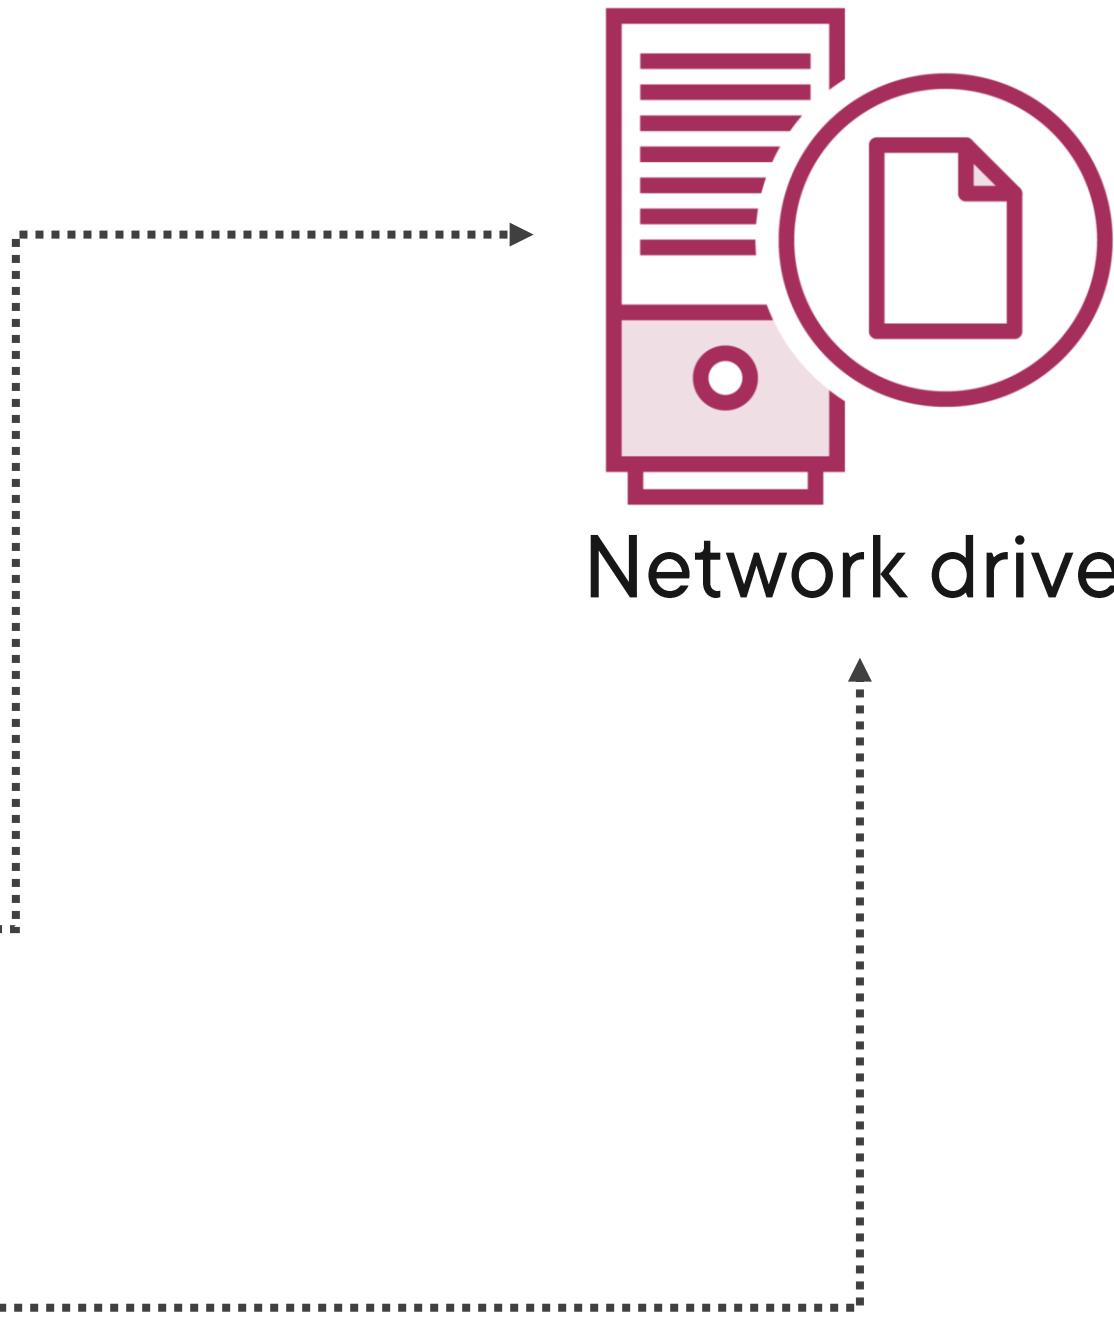




custom/path



another/custom/path



Changing MLflow Storage Location

```
mlflow.set_tracking_uri(  
    "file:///custom/location"  
)
```



in the Python code

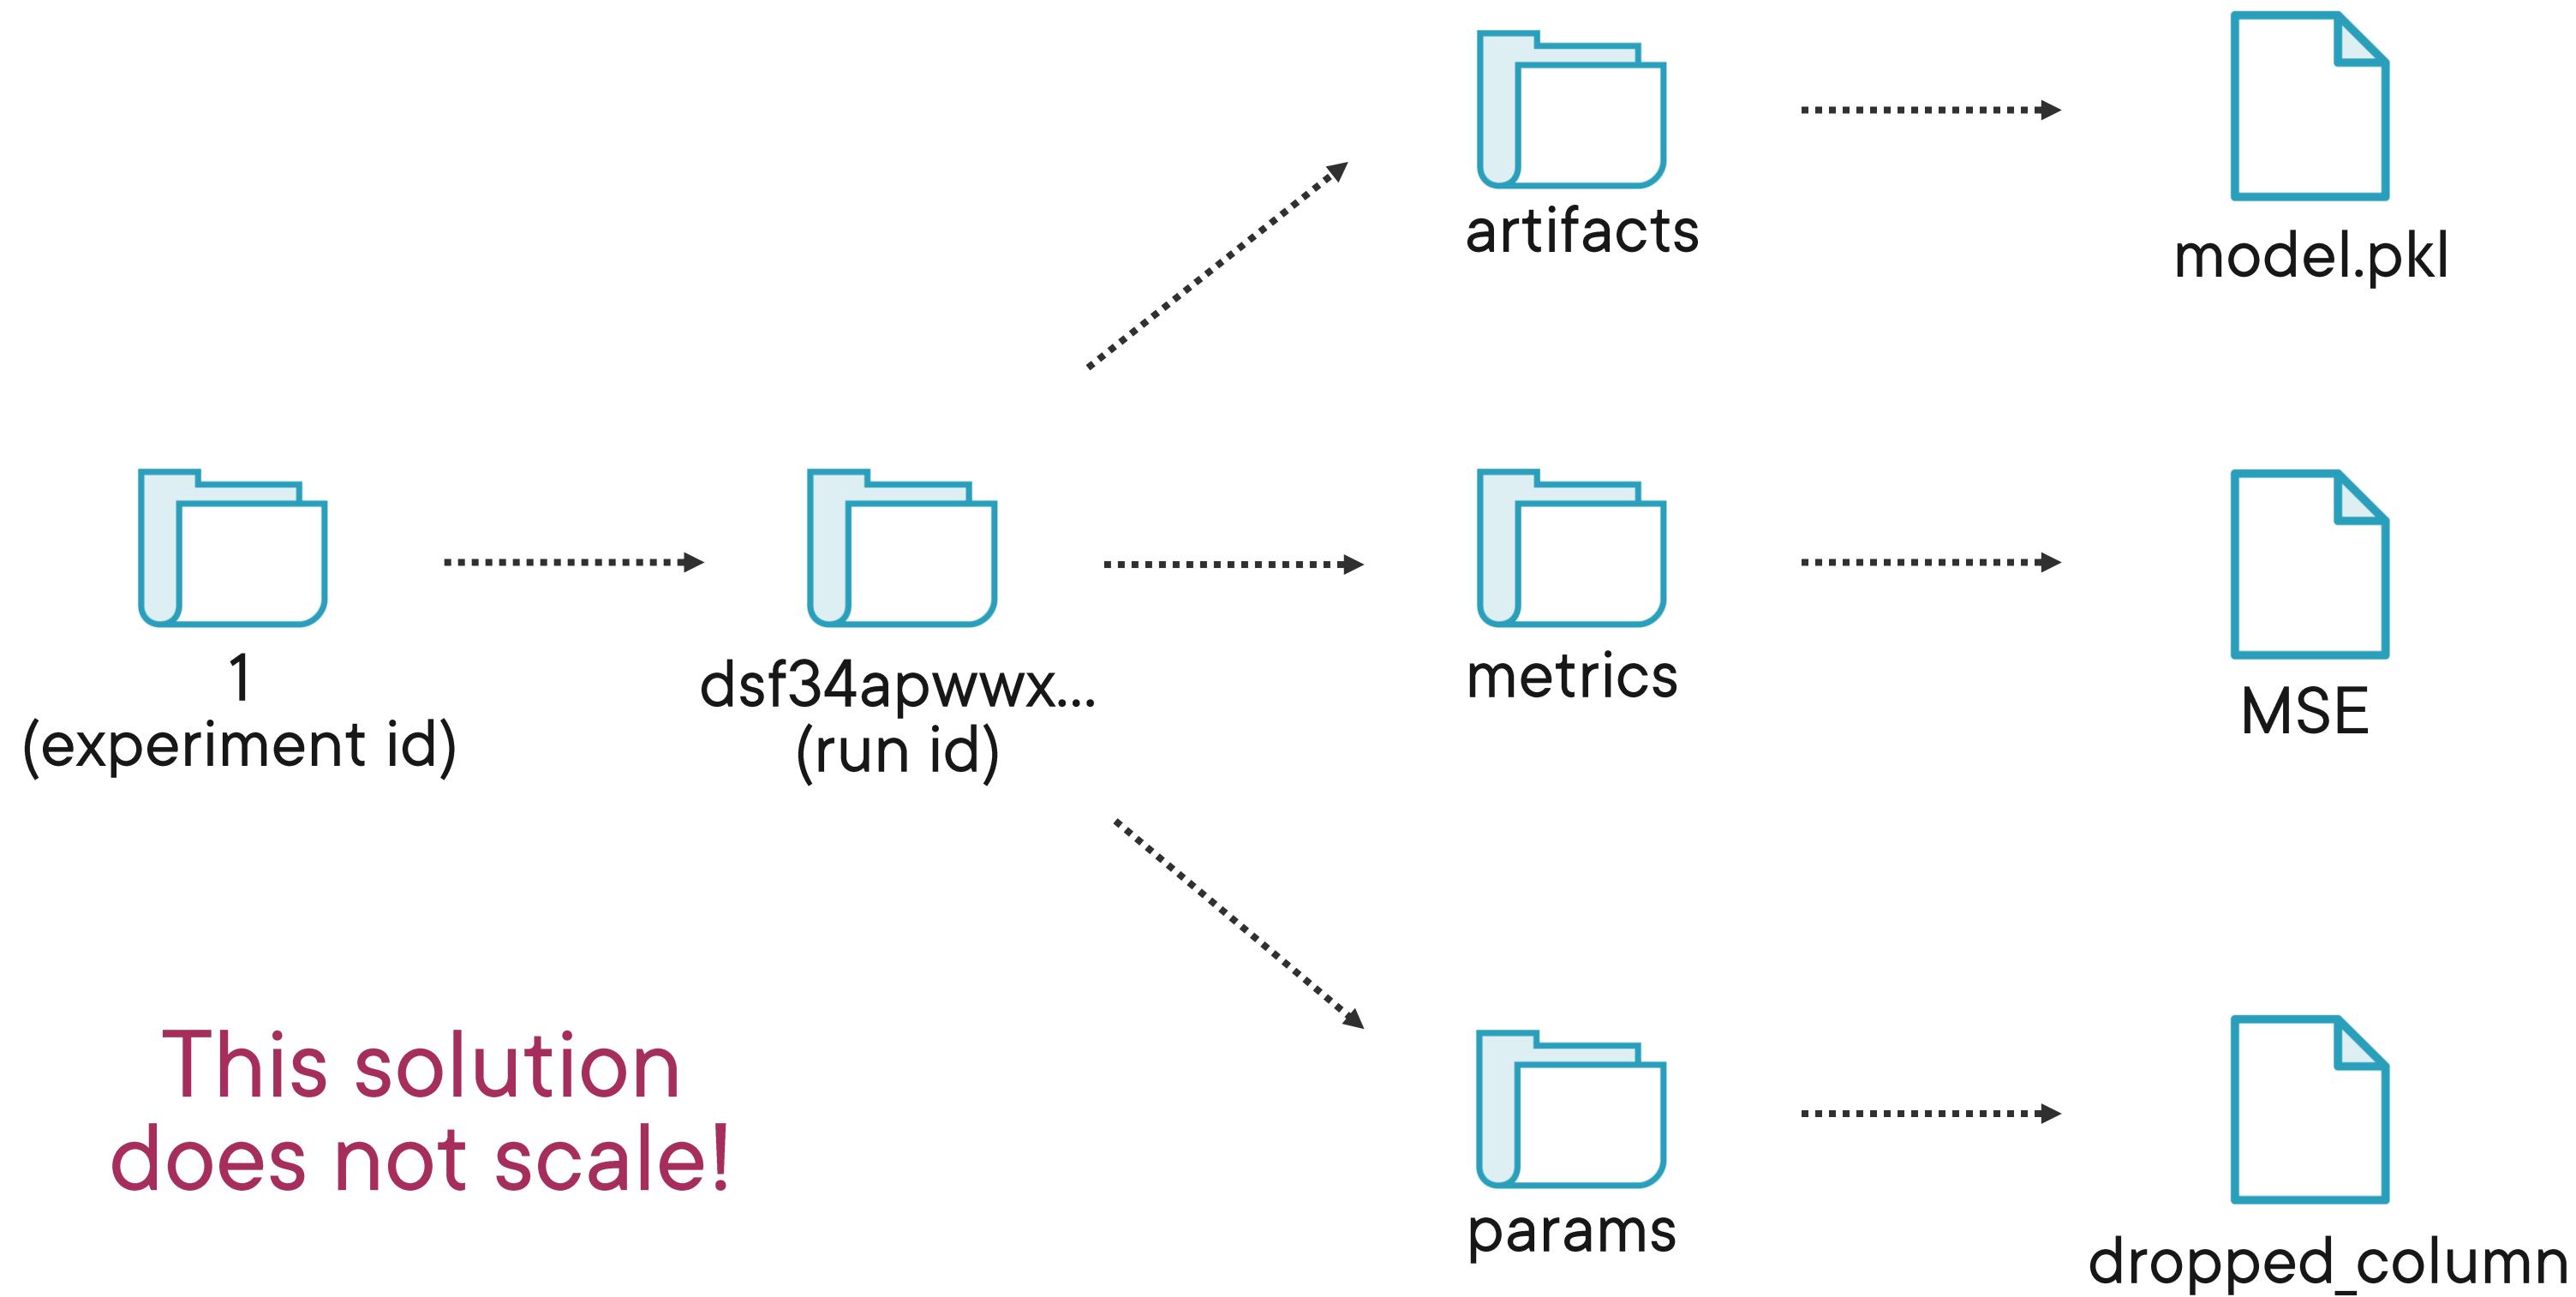
```
mlflow ui  
--backend-store-uri file:///custom/location
```



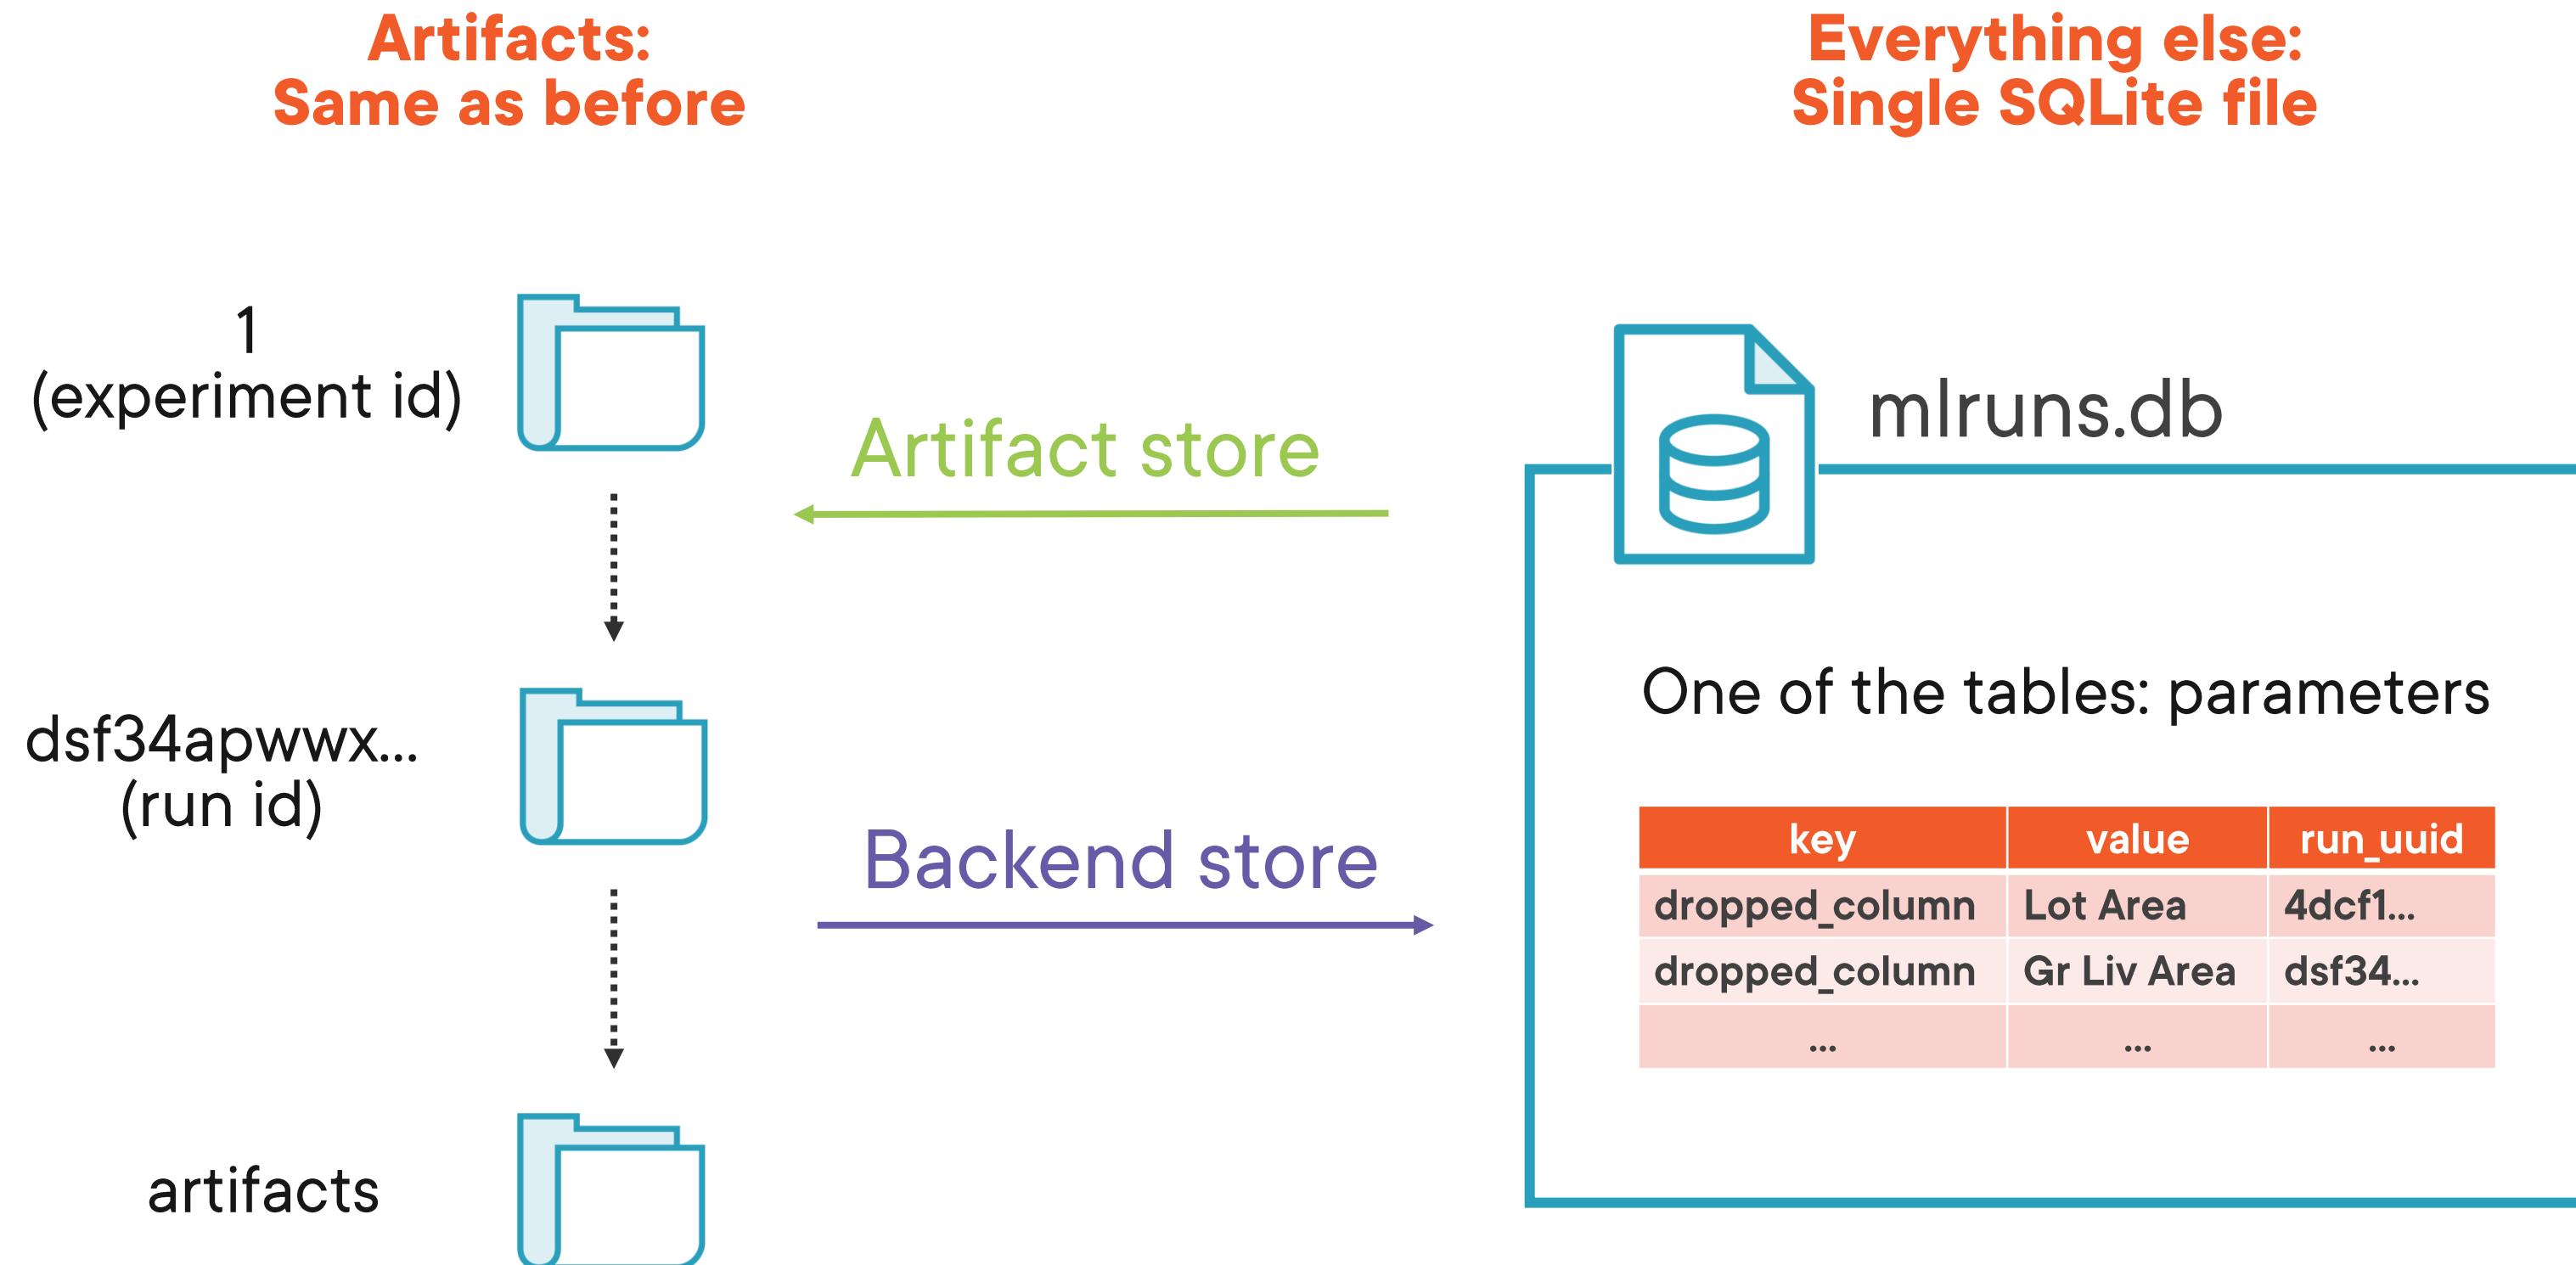
starting the UI



Inside 'mlruns'



Using SQLite



Using SQLite Backend Store

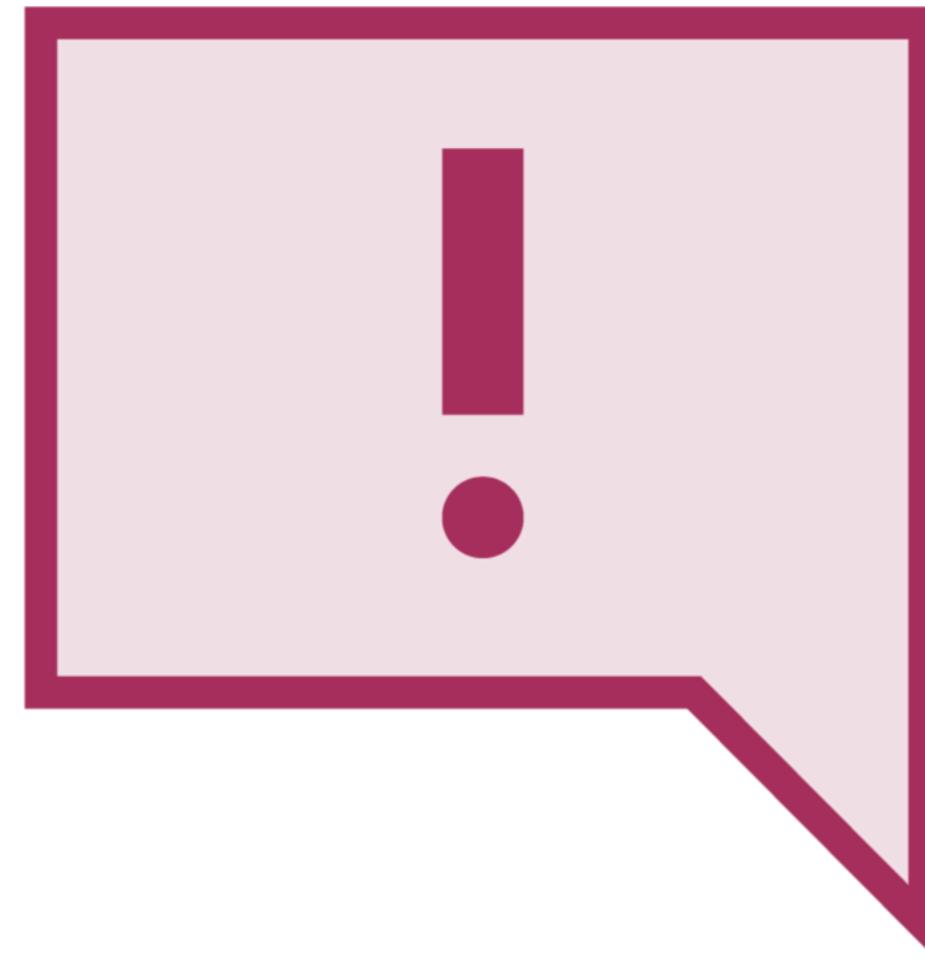
```
mlflow.set_tracking_uri(  
    "sqlite:///custom/location/mlruns.db"  
)
```

in the Python code

```
mlflow ui  
--backend-store-uri sqlite:///custom/location/mlruns.db
```

starting the UI





Artifact store cannot be customized in this setup.



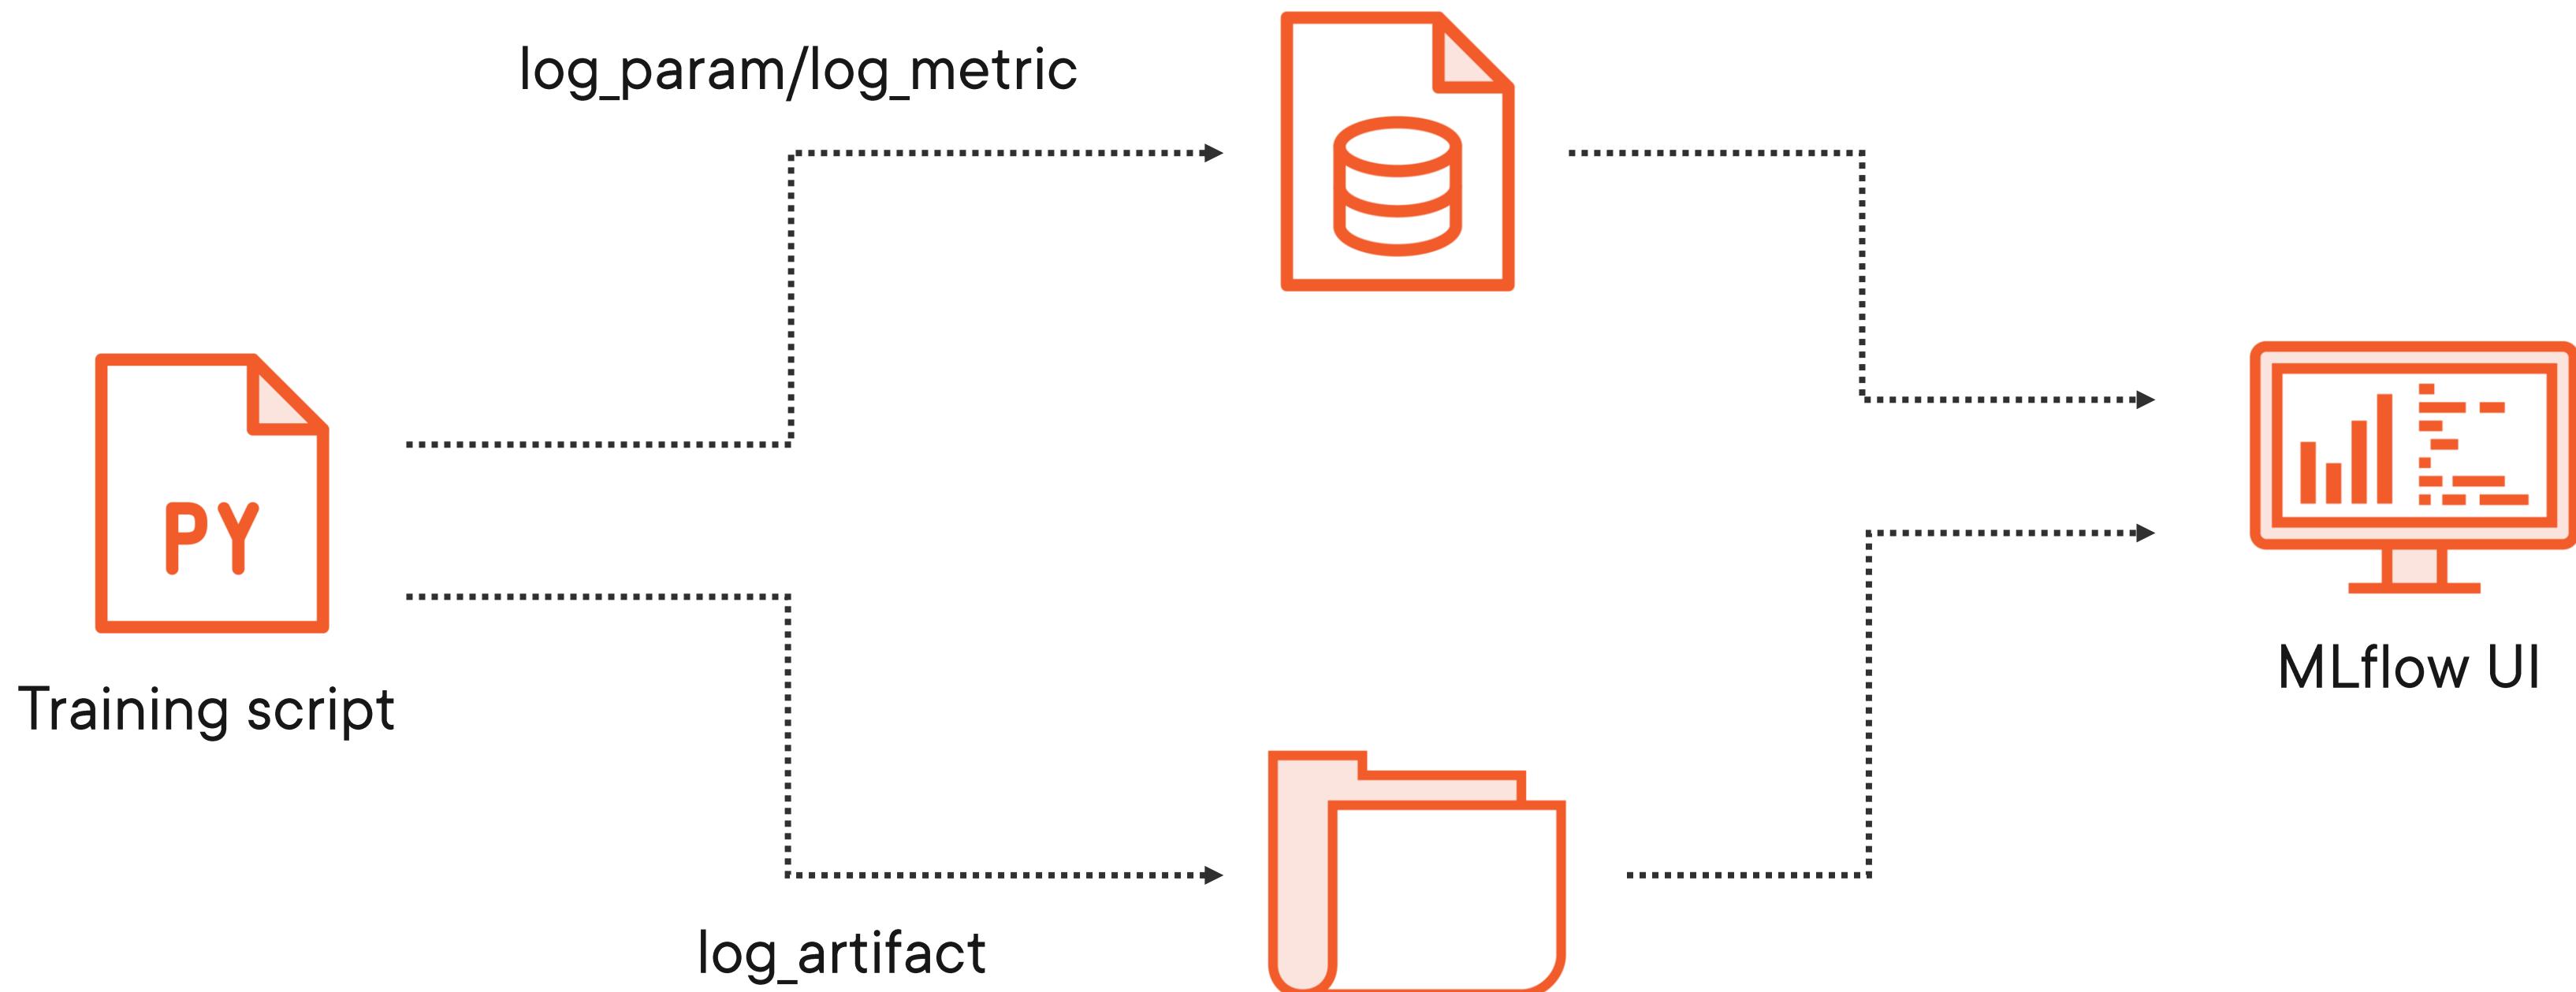
Up Next: MLflow Server



MLflow Server



Without MLflow Server



Using MLflow Server

```
mlflow server --host 0.0.0.0  
--backend-store-uri  
sqlite:///data/mlflow/mlruns.db  
--default-artifact-root  
file:///data/mlflow/artifacts
```



starting the server

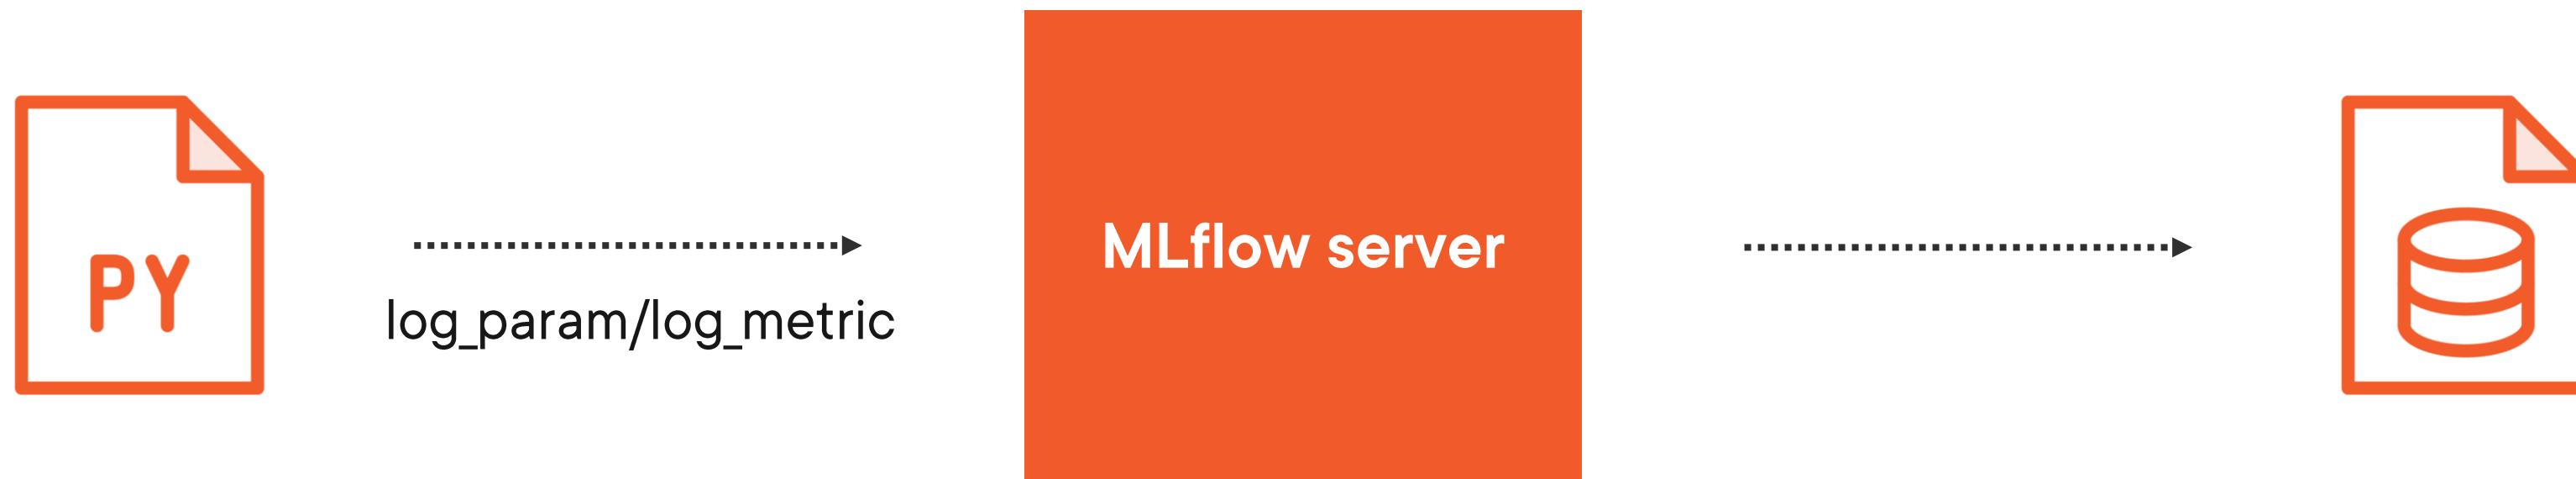
```
mlflow.set_tracking_uri("localhost:5000")
```



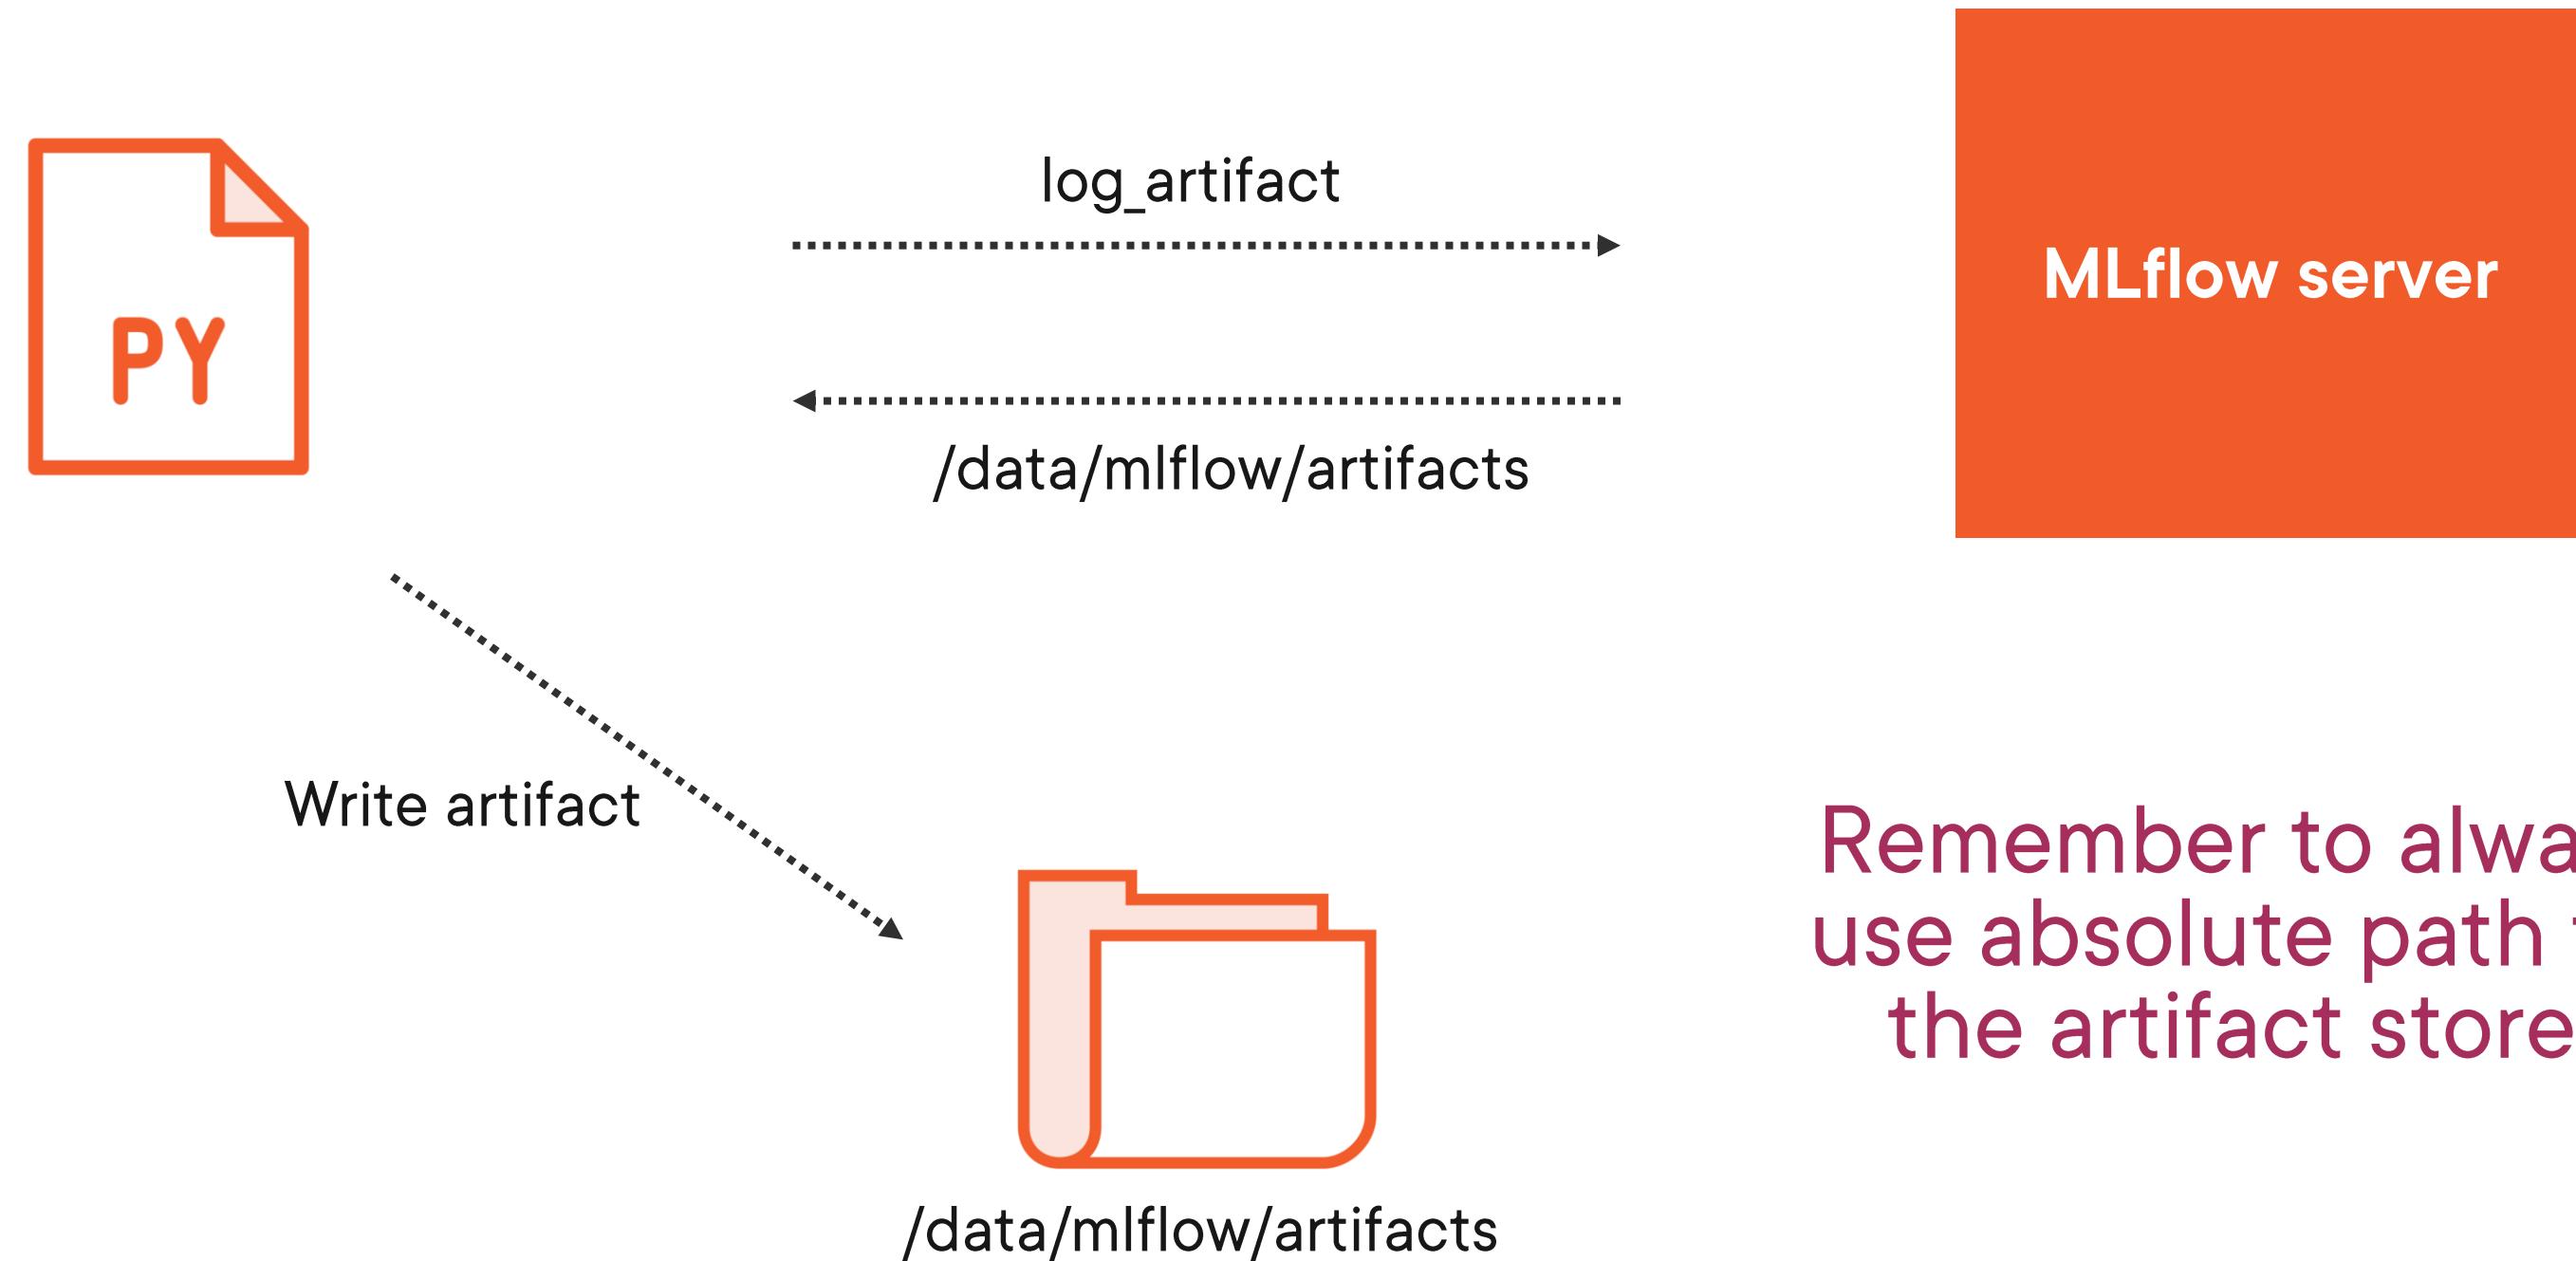
connecting to the
server in the Python
script



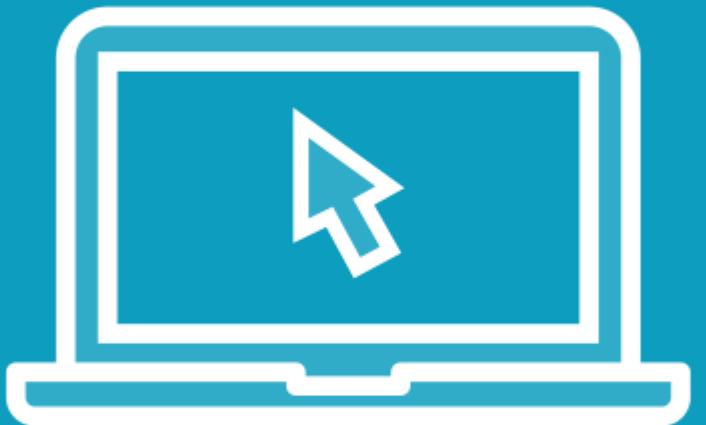
Backend Store with MLflow Server



Artifact Store with MLflow Server



Demo



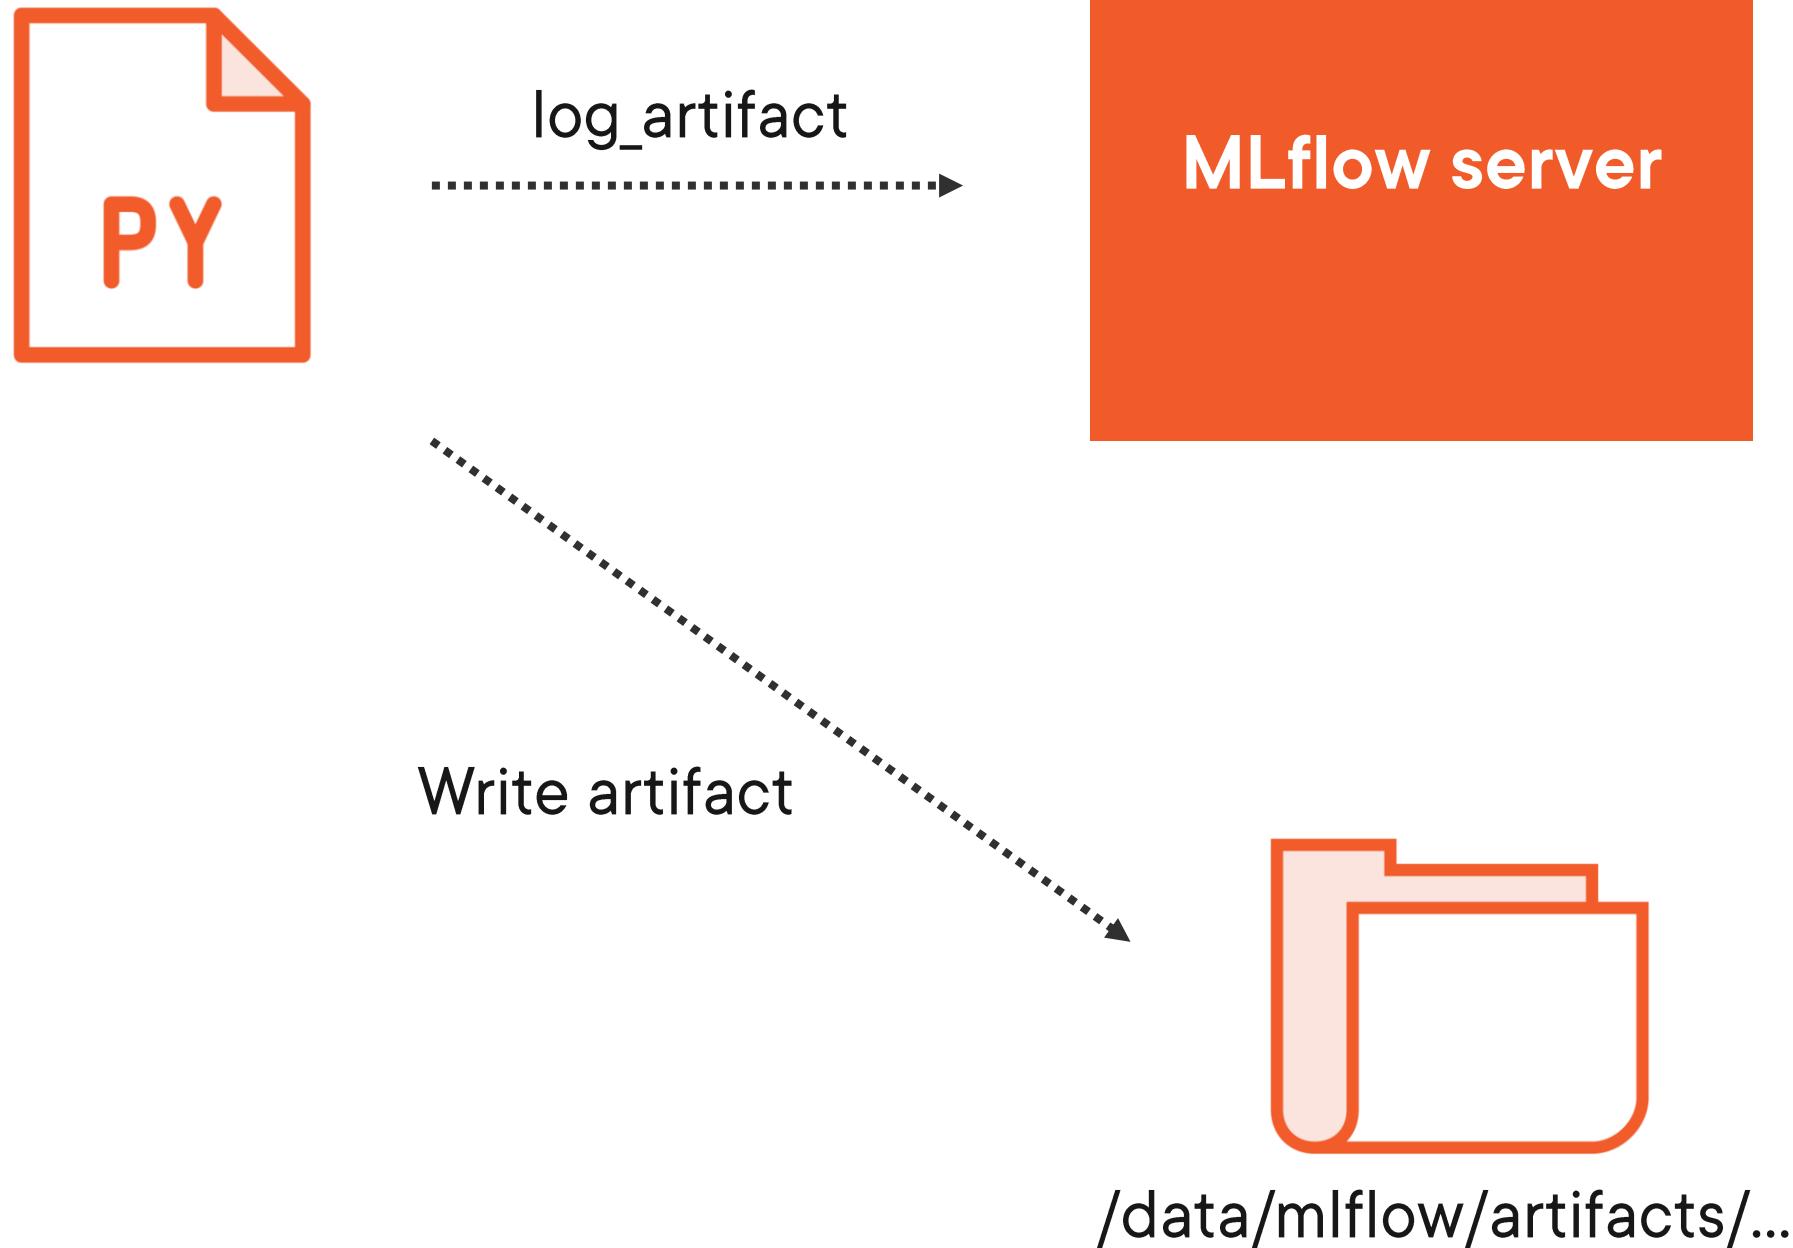
MLflow server



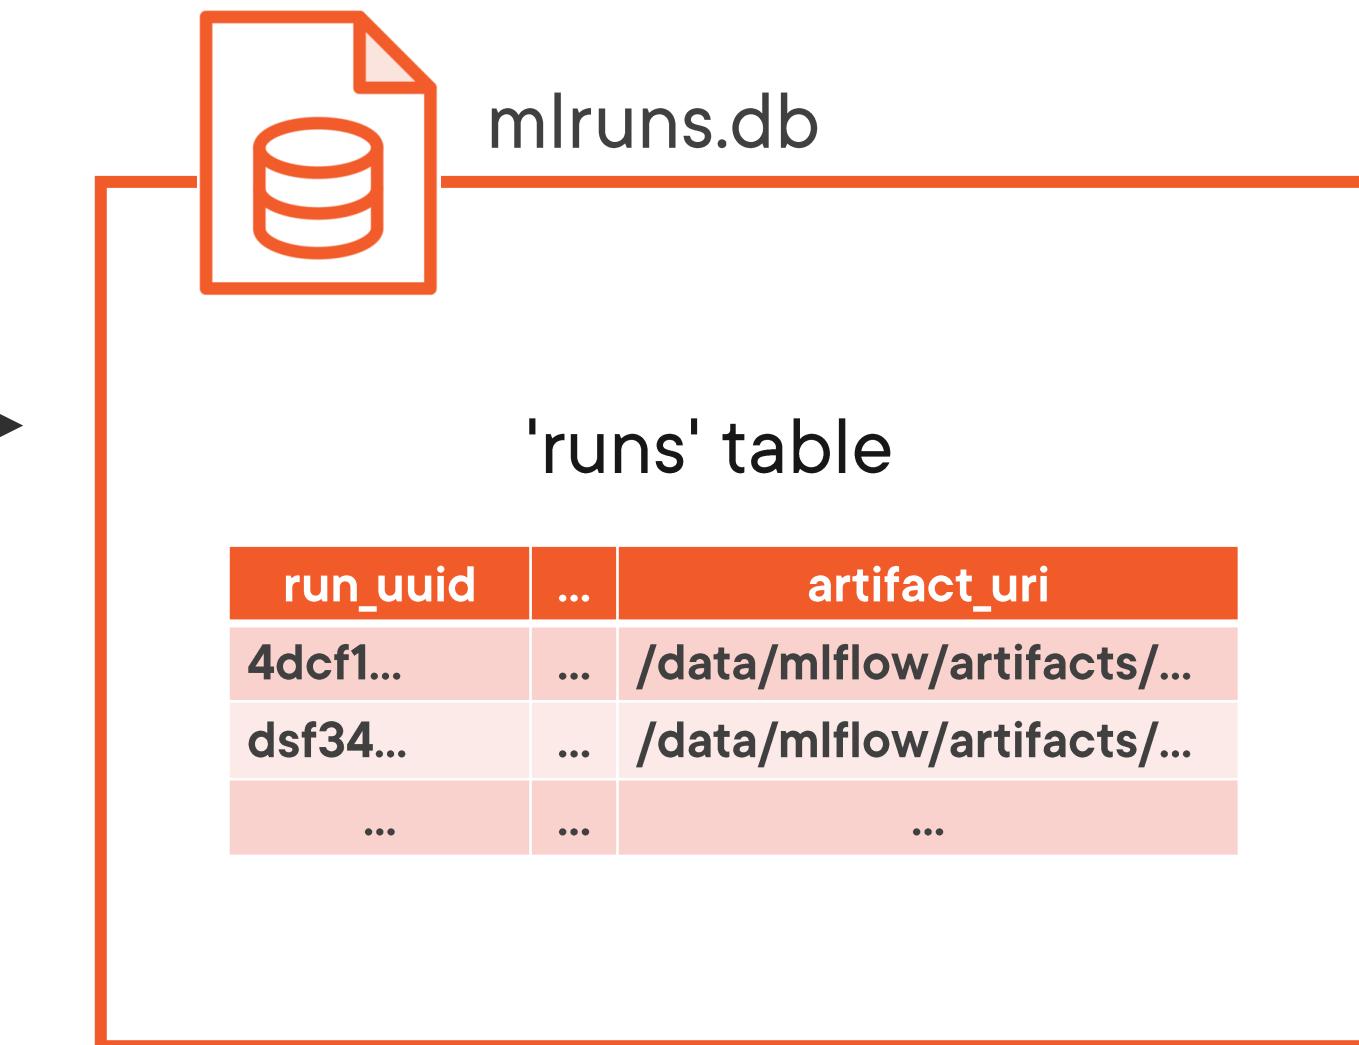
Complete MLflow Deployment



Everyone needs to run
their own MLflow server.



SQLite can be slow.



Every user needs to have the
artifacts folder mounted on the
same path.



Some Artifact Store Alternatives



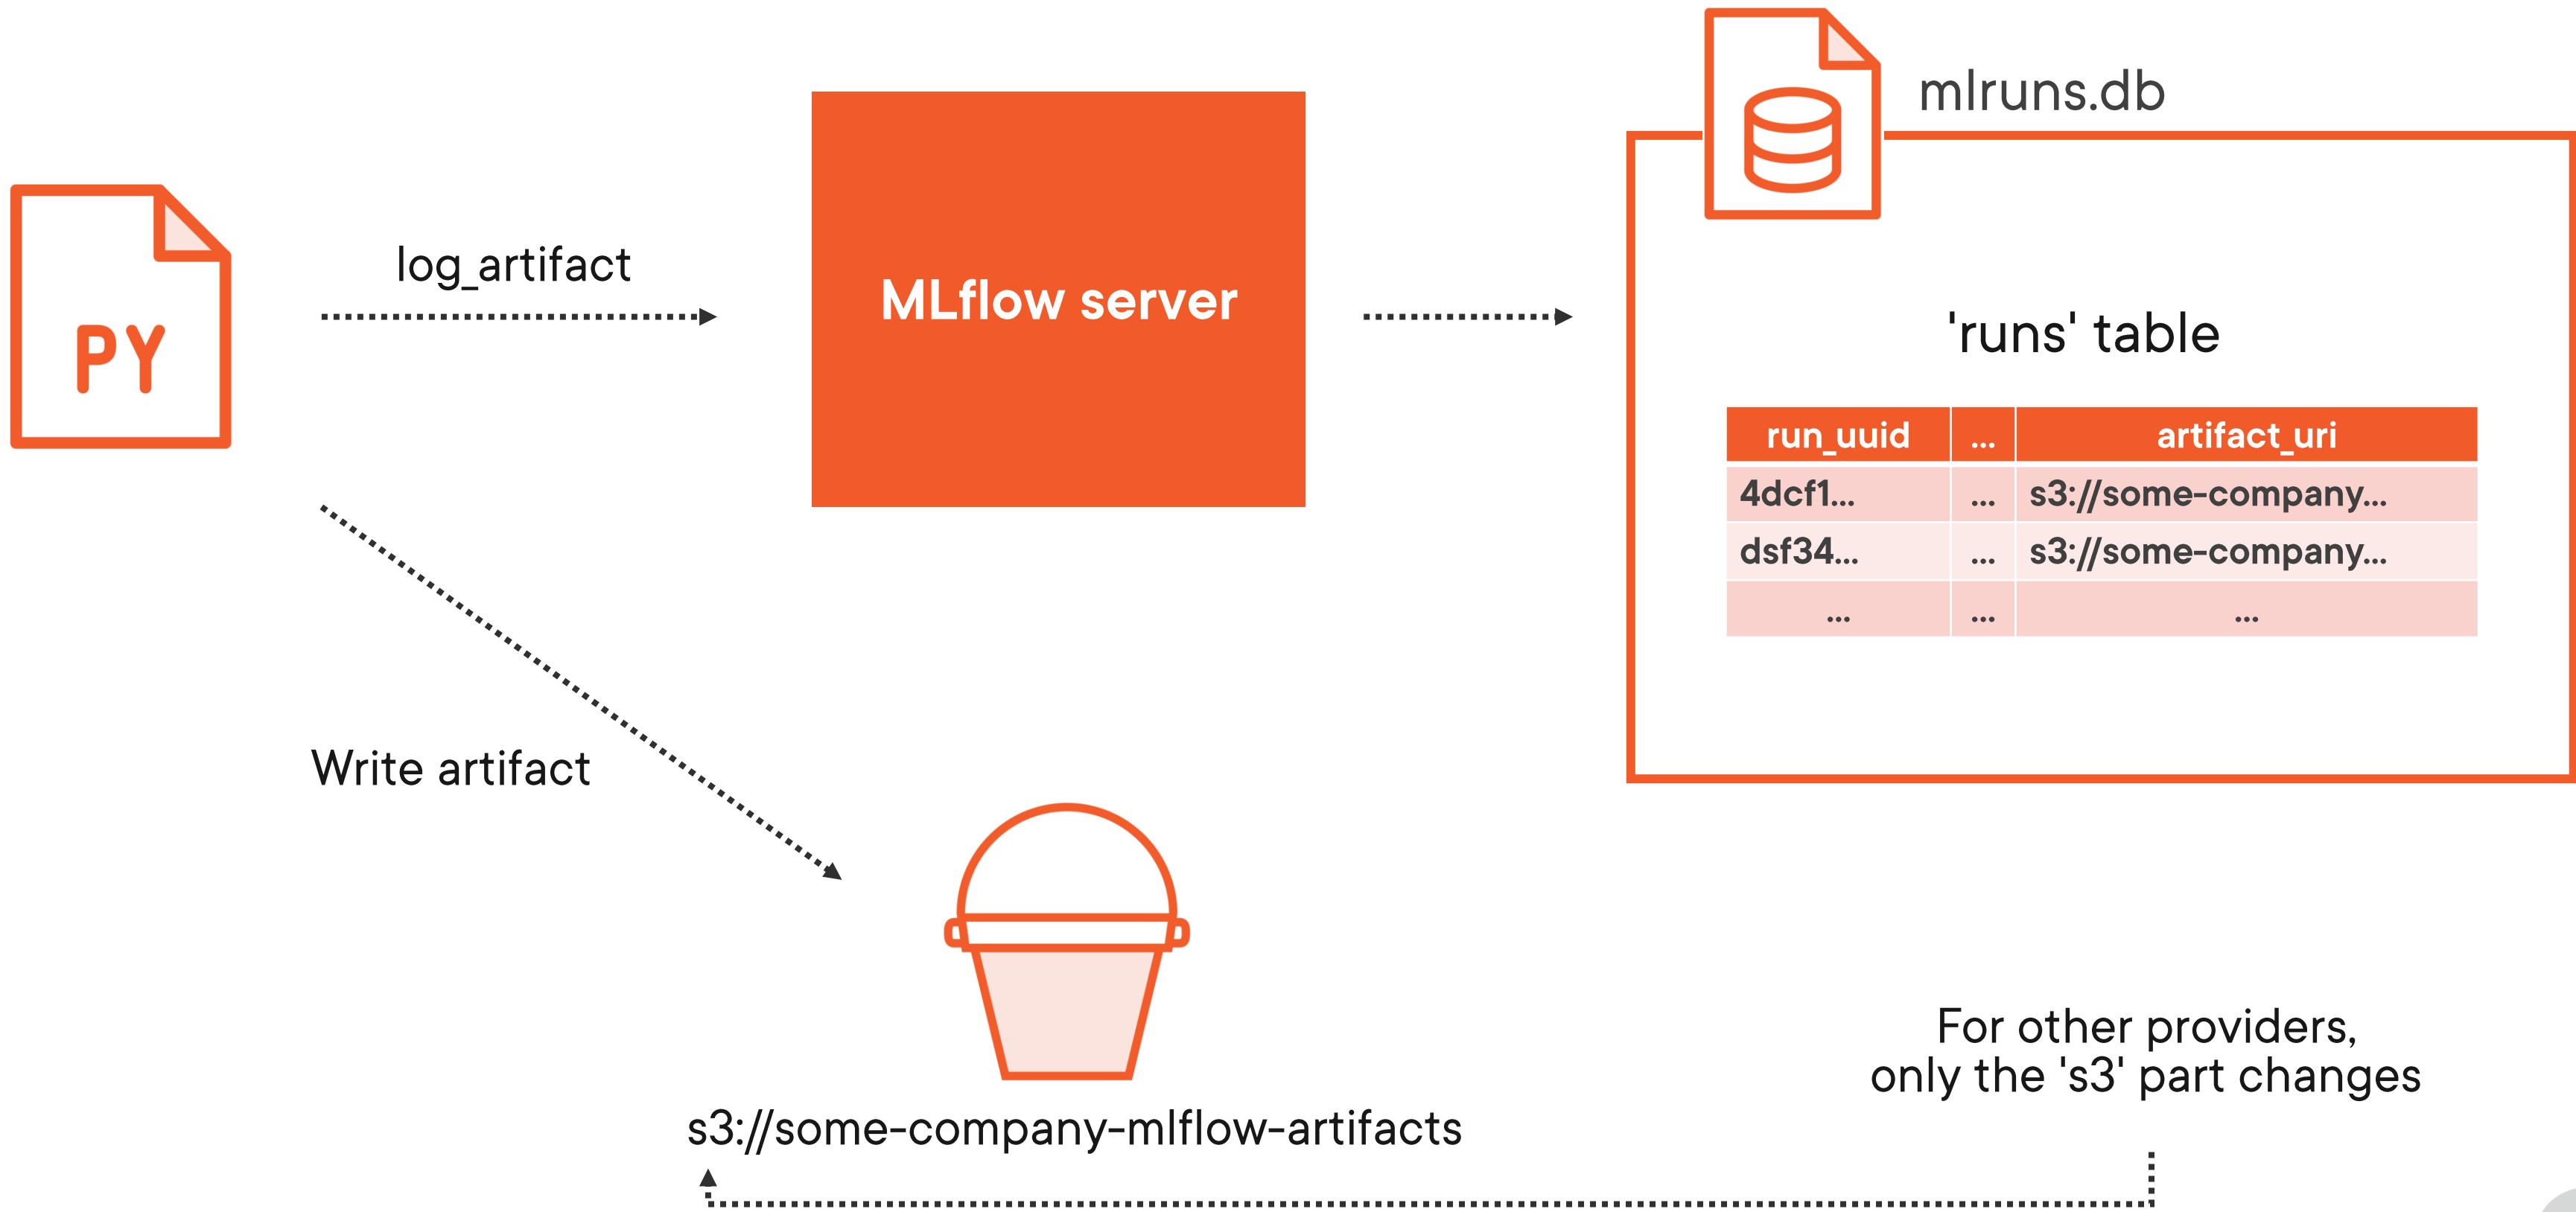
Some Artifact Store Alternatives

Amazon S3

Azure Blob Storage

**Google Cloud
Storage**





```
mlflow server --host 0.0.0.0  
--backend-store-uri  
sqlite:///data/mlflow/mlruns.db  
--default-artifact-root  
s3://some-company-mlflow-artifacts
```

Cloud Storage for Artifacts

Amazon S3

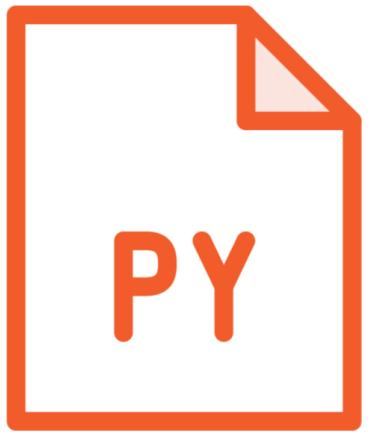
Other Artifact Store Options

(S)FTP

HDFS



Remote Server



```
mlflow.set_tracking_uri(  
    "mlflow.companyname.net"  
)
```



mlflow.companyname.net



Backend store



Summary



Shortcomings of the default setup

Using SQLite

Distinction between artifact store and backend store

MLflow server

Architecture of a complete MLflow deployment



Packaging and Running Models



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Overview



Models as a special type of artifact

Models directly supported by MLflow

Custom models

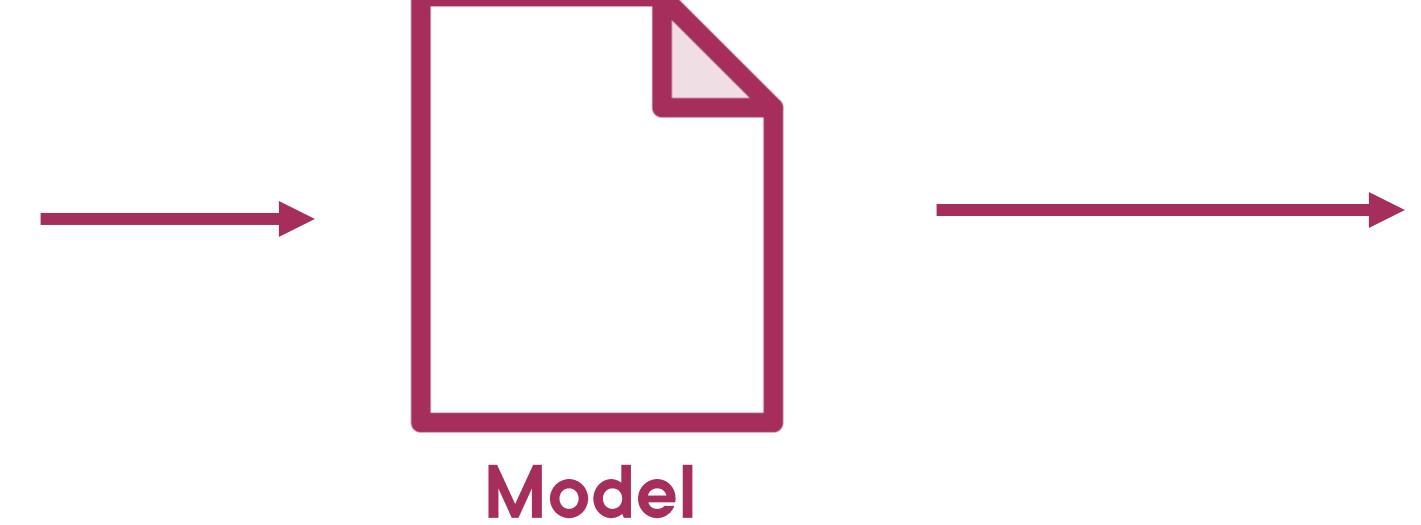
Running predictions



Logging the Model



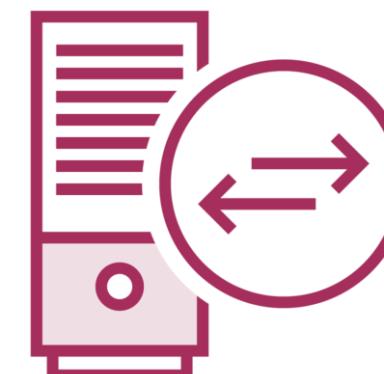
Training and evaluation



Production environment



Batch predictions



Real-time predictions

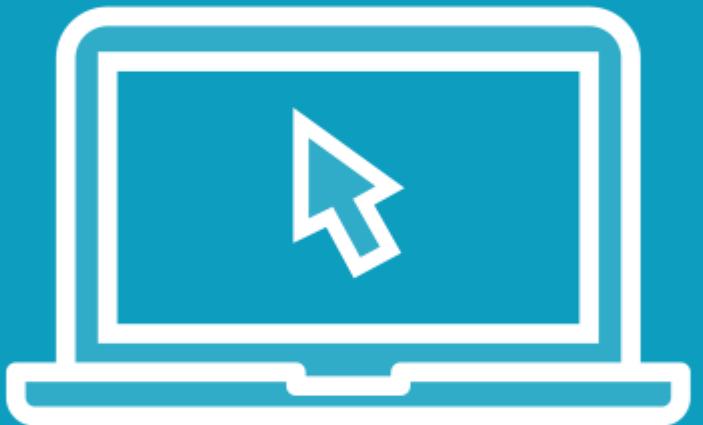


```
lr_model = LinearRegression()  
lr_model.fit(...)  
  
mlflow.sklearn.log_model(lr_model, "model")
```

Logging a Model

Supported 'flavor' – e.g. scikit-learn

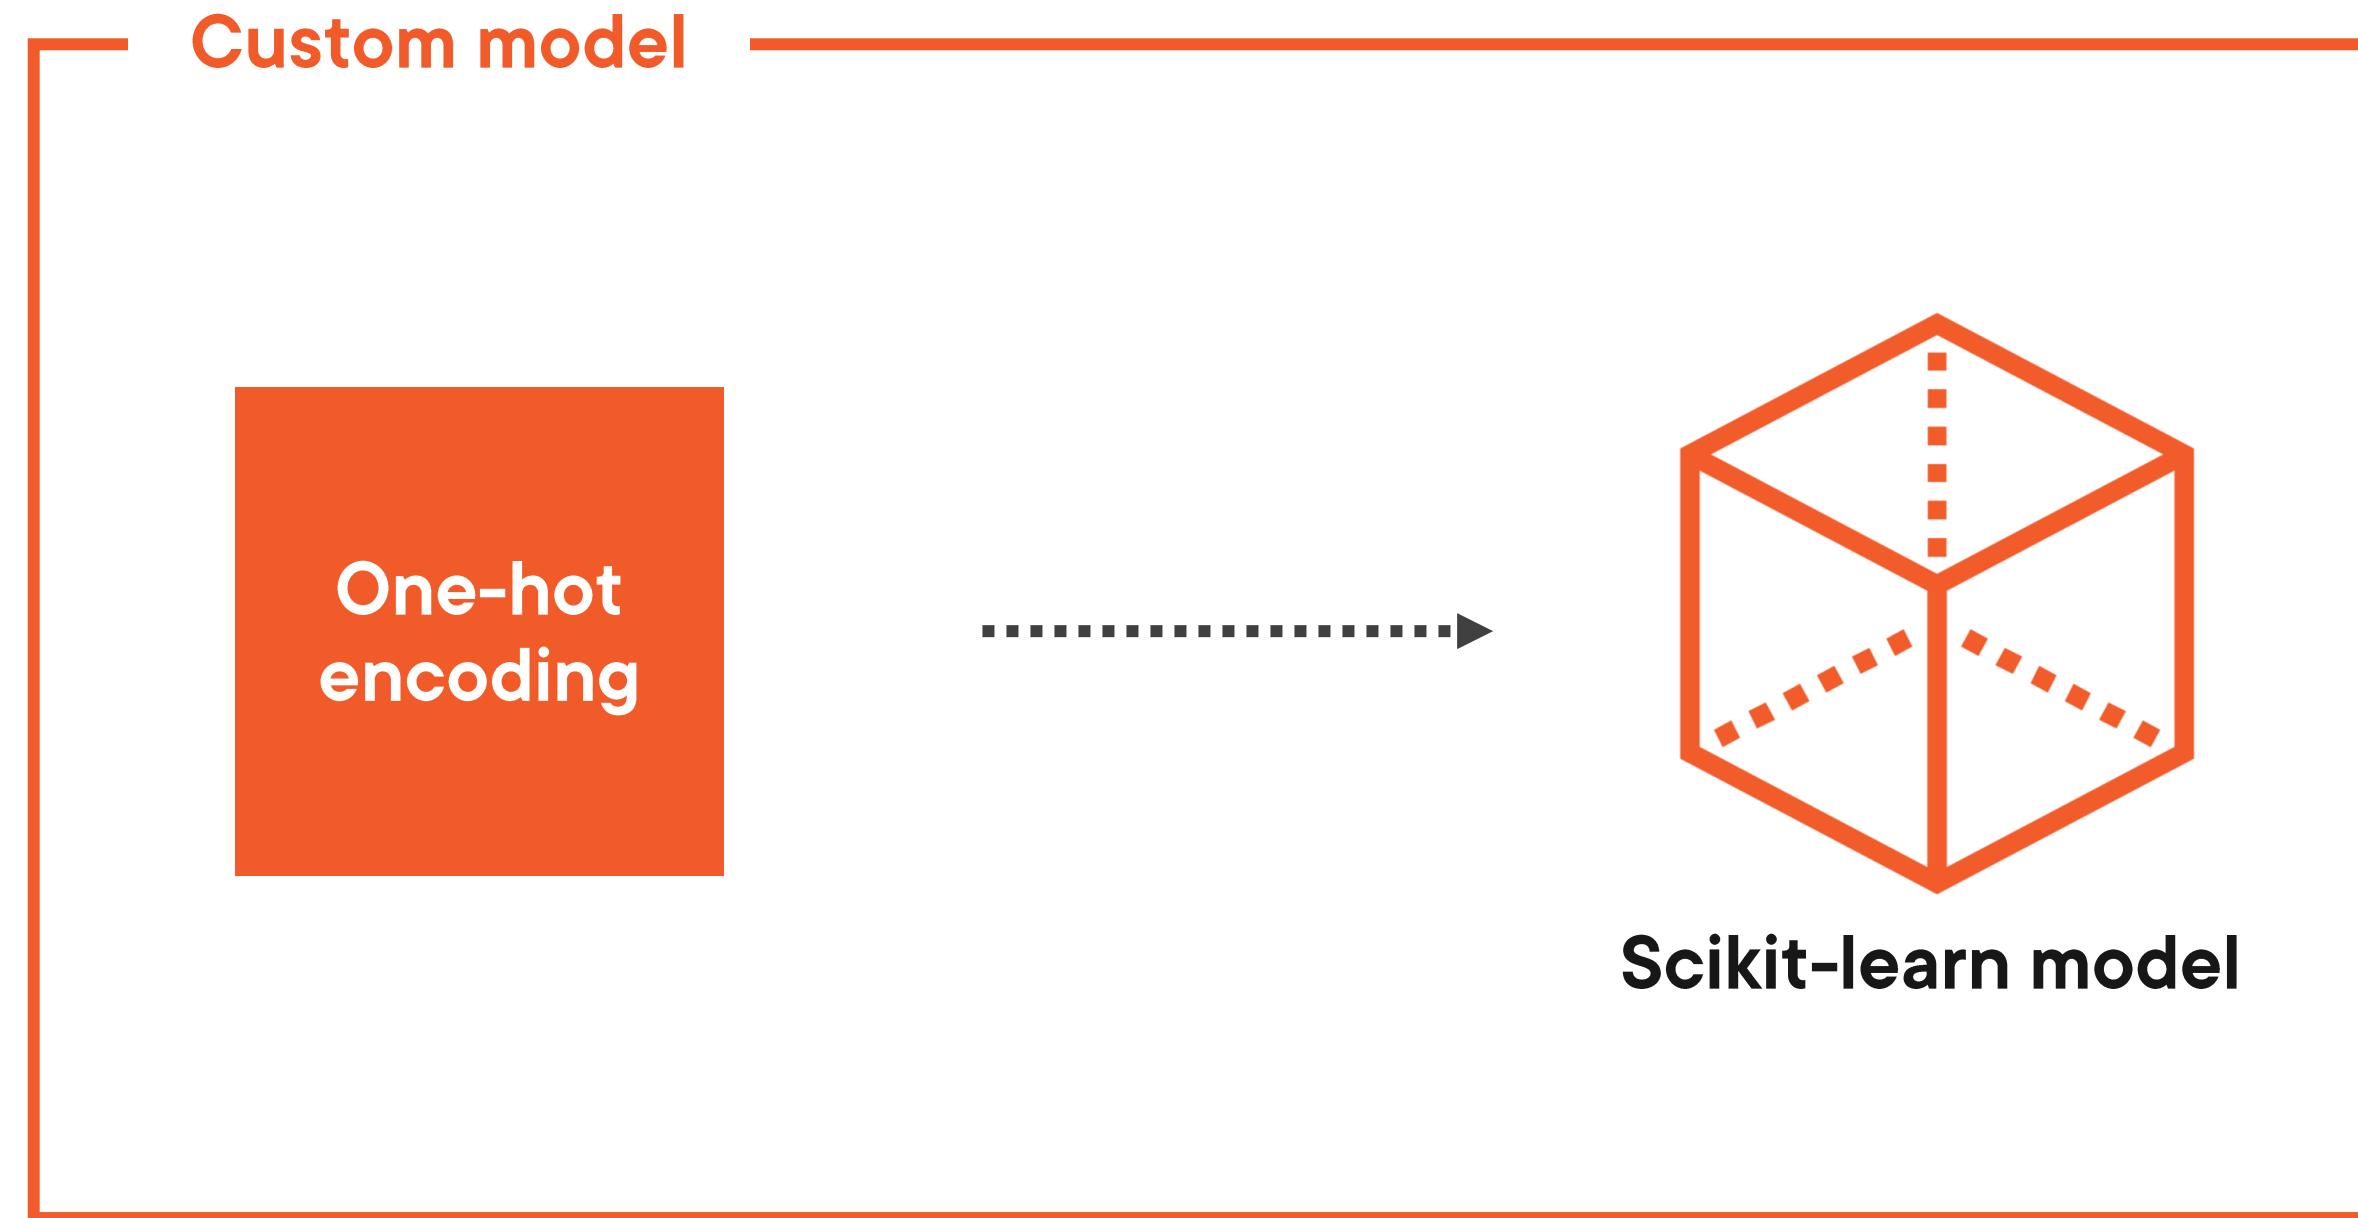
Demo



Logging the trained model



Customizing the Model



Creating a Custom Model

'pyfunc' flavor

```
from mlflow.pyfunc import PythonModel

class CustomModel(PythonModel):

    def predict(self, ..., model_input):
        return np.array(...)

mlflow.pyfunc.log_model("model", python_model=CustomModel())
```



Storing the Original Model

```
sklearn_model = LinearRegression()  
... # Train the model  
  
with open("tmp/model.pkl", "wb") as m:  
    m.write(sklearn_model)  
  
artifacts = { "sklearn_model": "tmp/model.pkl" }  
  
mlflow.pyfunc.log_model("model", python_model=CustomModel(), artifacts=artifacts)
```



Loading the Original Model

```
from mlflow.pyfunc import PythonModel

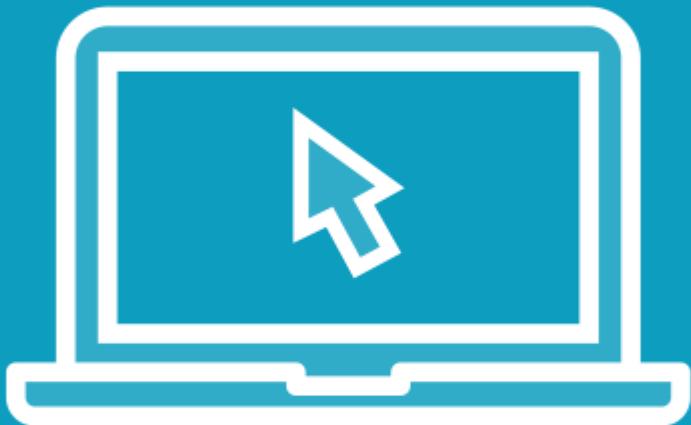
class CustomModel(PythonModel):

    def load_context(self, context):
        with open(context.artifacts["sklearn_model"]) as m:
            self.sklearn_model = pickle.load(m)

    def predict(self, ..., model_input):
        return self.sklearn_model.predict(model_input)
```



Demo



Custom model class
One-hot encoding



Running Predictions



Batch Predictions

```
mlflow models predict  
-m <MODEL_PATH>  
-i <INPUT_DATASET_FILE>  
--no-conda
```



Batch Predictions

```
mlflow models predict  
-m runs:/de223jj8sd.../model  
-i <INPUT_DATASET_FILE>  
--no-conda
```

Run ID

Artifacts subdirectory for the model



Batch Predictions

```
$Env:MLFLOW_TRACKING_URI = "http://localhost:5000" ←..... Windows  
export MLFLOW_TRACKING_URI="http://localhost:5000" ←..... Linux/macOS
```

```
mlflow models predict  
-m runs:/de223jj8sd.../model  
-i <INPUT_DATASET_FILE>  
--no-conda
```



Batch Predictions

```
$Env:MLFLOW_TRACKING_URI = "http://localhost:5000"  
export MLFLOW_TRACKING_URI="http://localhost:5000"  
  
mlflow models predict  
-m runs:/de223jj8sd.../model  
-i <INPUT_DATASET_FILE> ←----- CSV or JSON – Pandas compatible  
--no-conda
```



Batch Predictions

```
$Env:MLFLOW_TRACKING_URI = "http://localhost:5000"  
export MLFLOW_TRACKING_URI="http://localhost:5000"
```

```
mlflow models predict  
-m runs:/de223jj8sd.../model  
-i <INPUT_DATASET_FILE>  
--no-conda  
-t csv
```

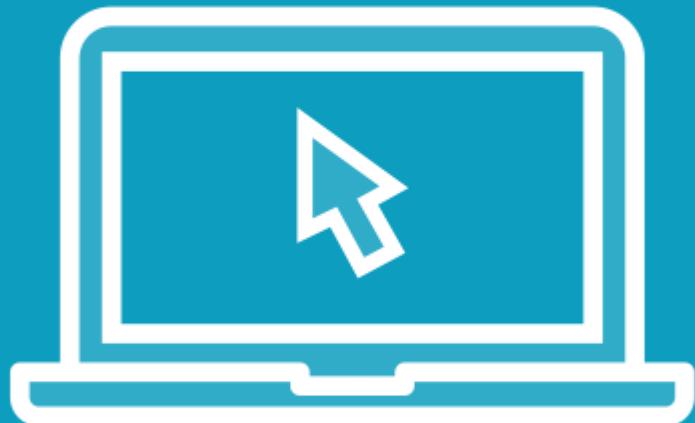


Batch Predictions

```
$Env:MLFLOW_TRACKING_URI = "http://localhost:5000"  
export MLFLOW_TRACKING_URI="http://localhost:5000"  
  
mlflow models predict  
-m runs:/de223jj8sd.../model  
-i <INPUT_DATASET_FILE>  
--no-conda  
-t csv  
-o results.json
```



Demo



Predictions using existing model

Static dataset

Real-time predictions



Summary



Models as a special class of artifacts

Custom models

Running predictions



Sharing and Managing Models with Model Registry



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Overview



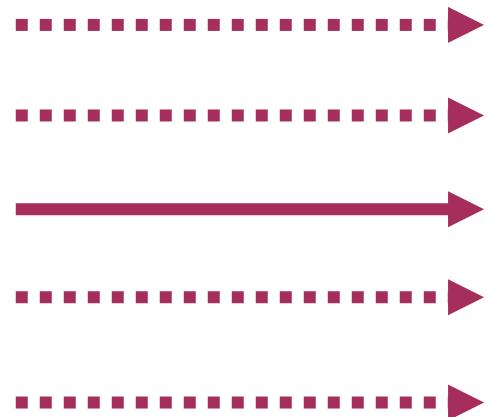
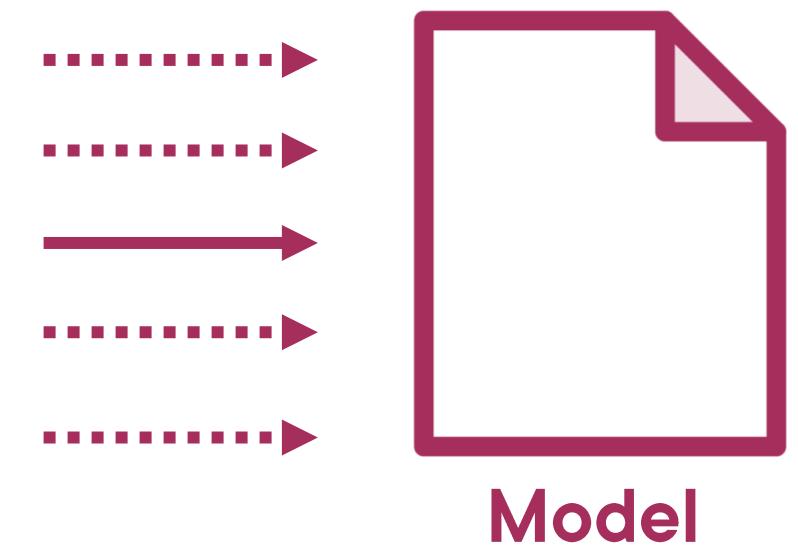
Problems that Model Registry solves

Working with Model Registry in the UI

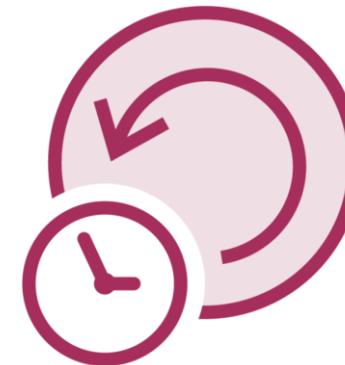
Using Model Registry when training and running predictions



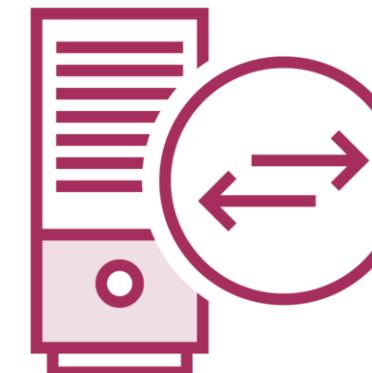
Training and evaluation



Production environment



Batch
predictions



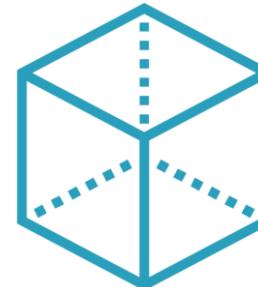
Real-time
predictions



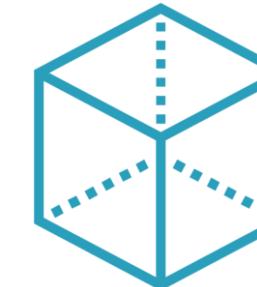
Production environment



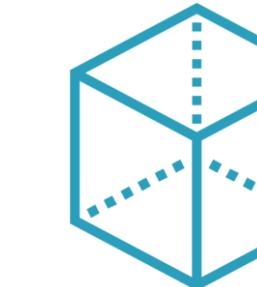
Batch
predictions



House prices

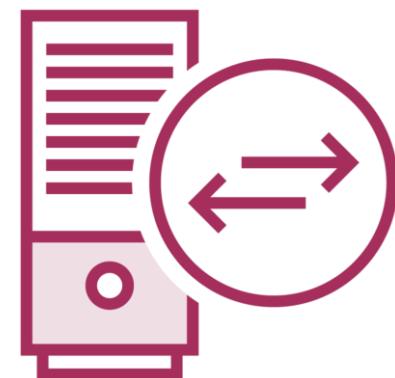


Focus areas

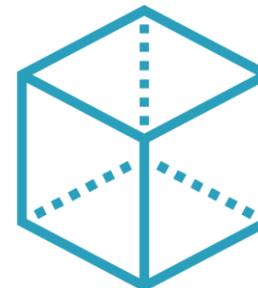


Another task

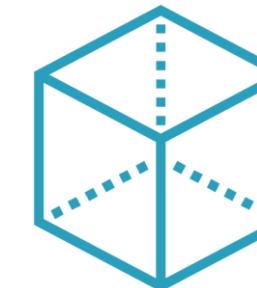
...



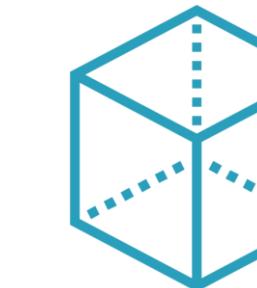
Real-time
predictions



House prices



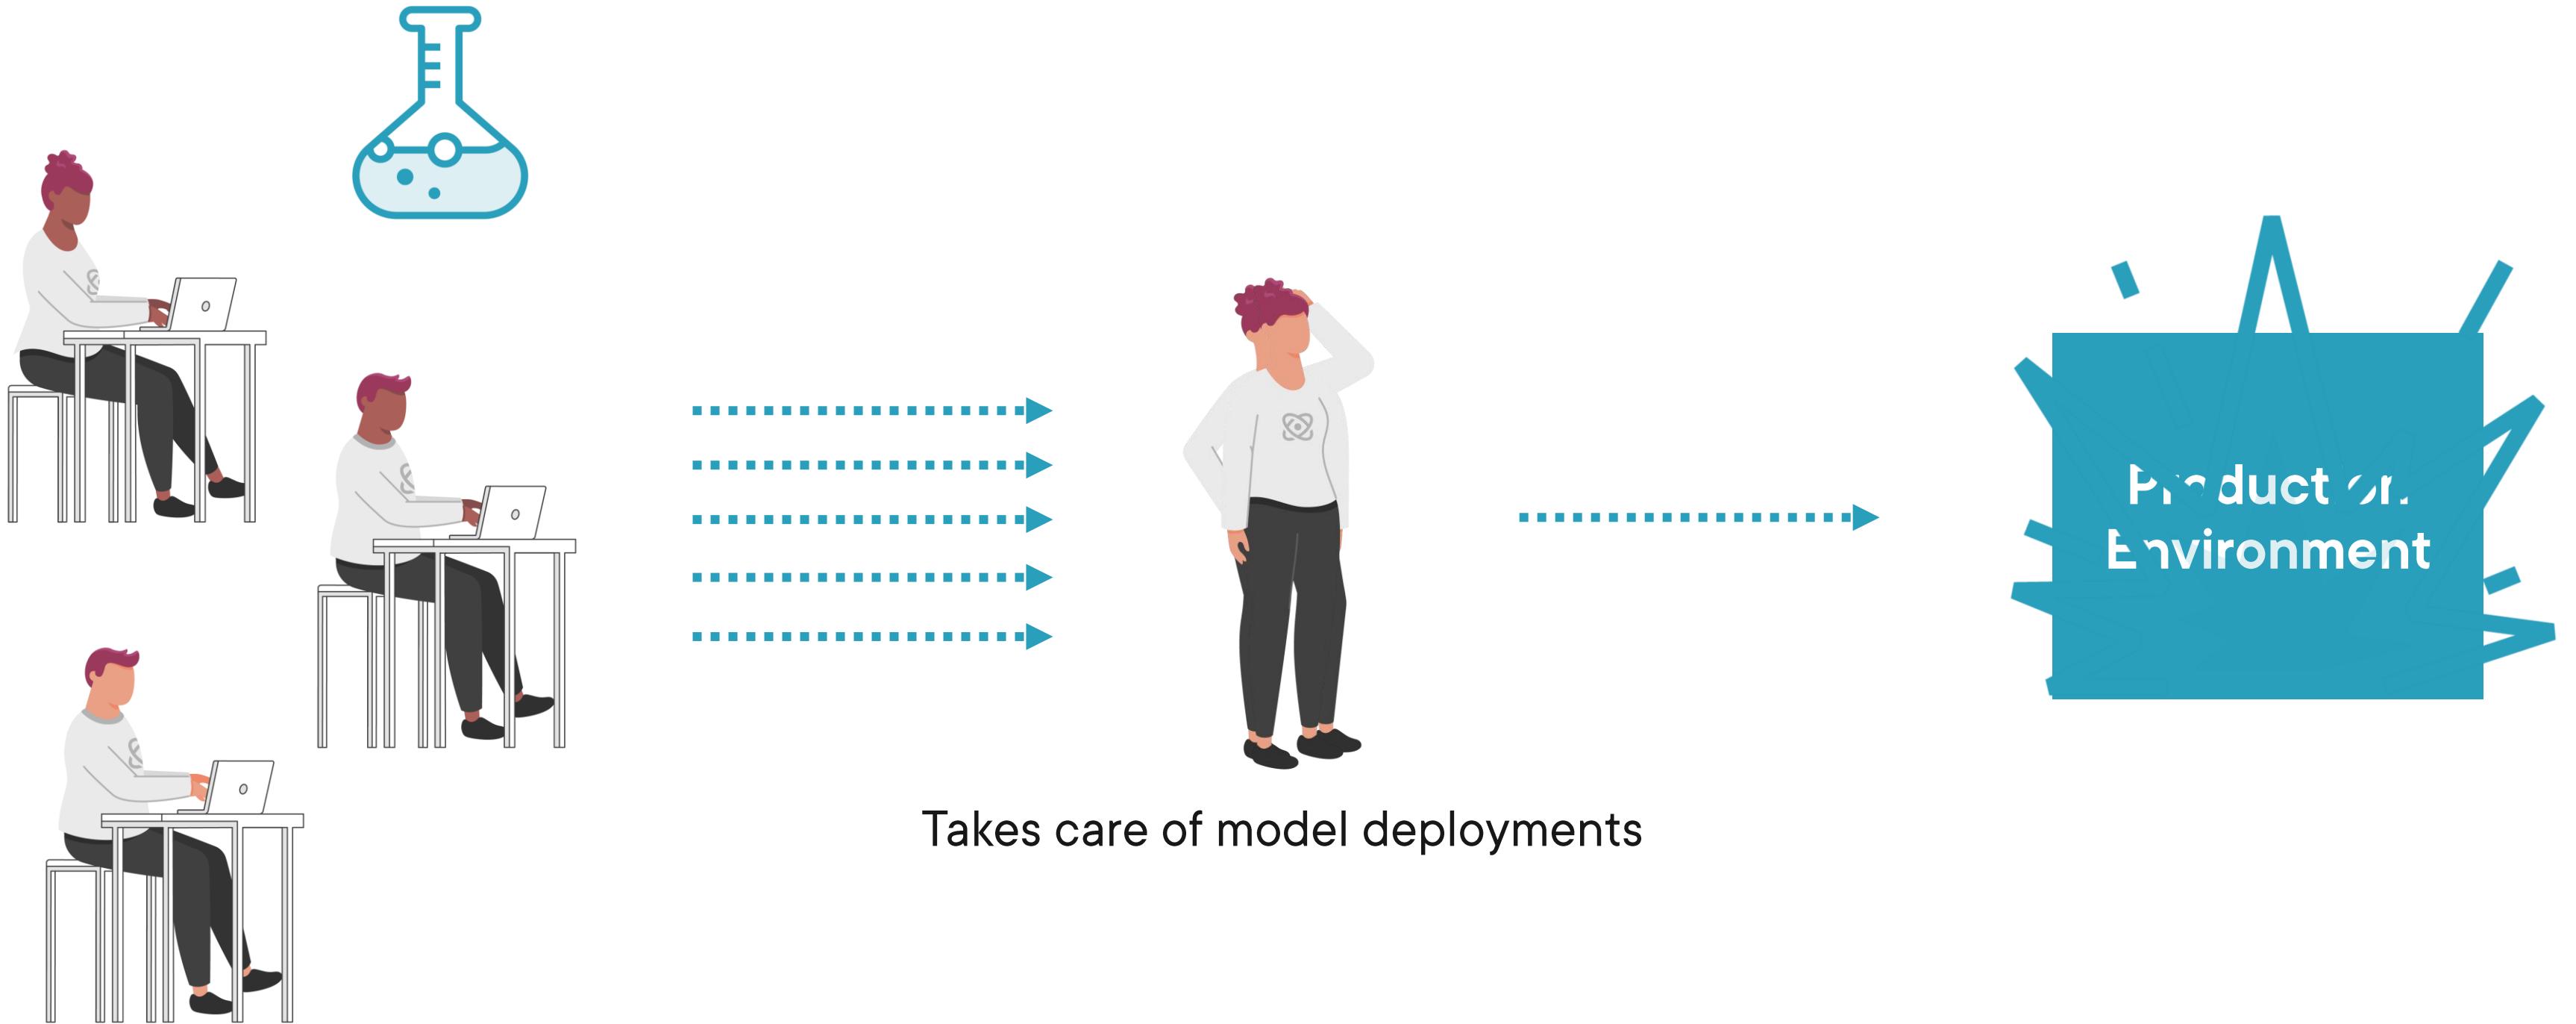
Focus areas



Another task

...

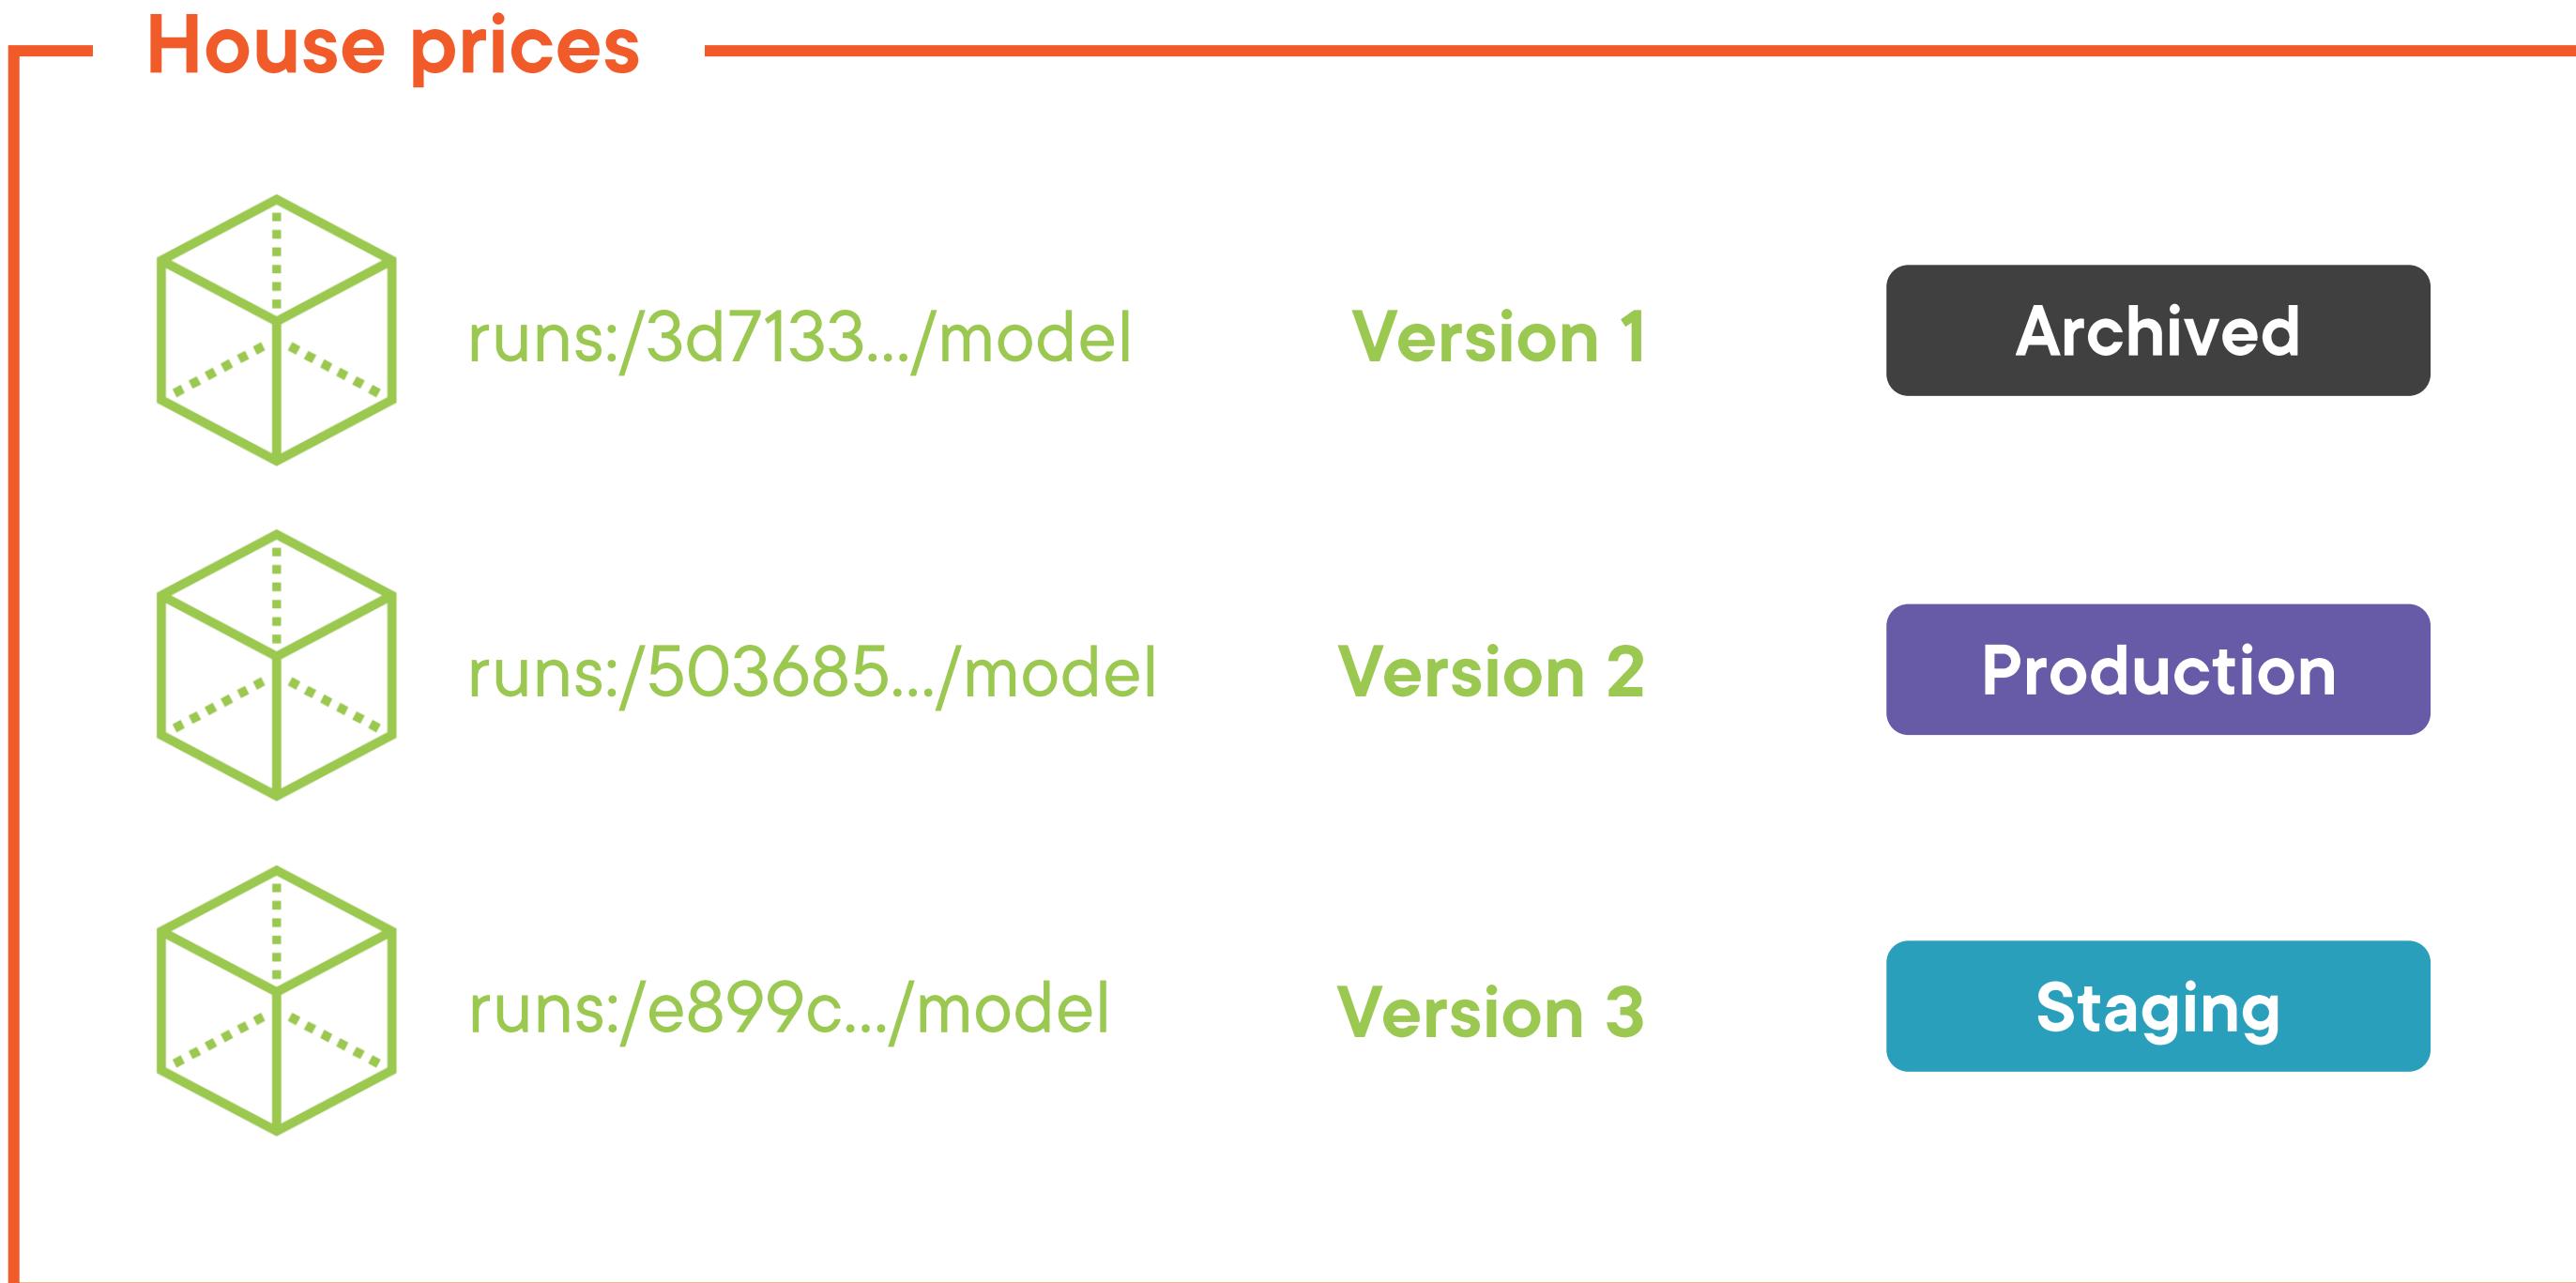




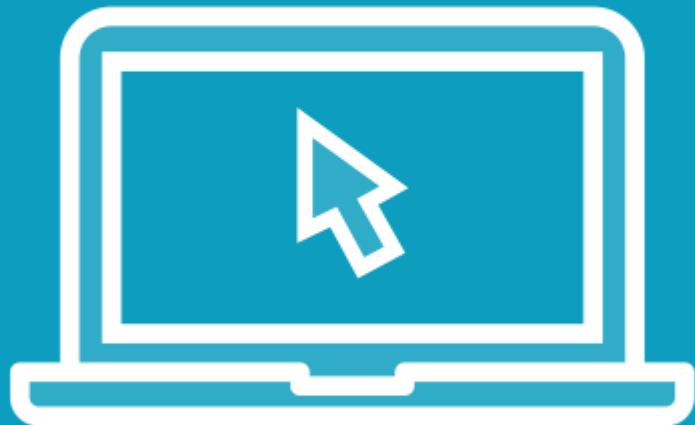
Model Registry



Registered Model



Demo



Creating registered models

Registering individual models

Managing stages

Running predictions



Summary



Problems that Model Registry solves

Registered models, versions, and stages

Registering models from the UI

Registering models from training scripts

Using models from the registry



Summary



Paweł Kordek

Software Engineer – Data and Machine Learning

@pawel_kordek



Machine Learning Workflows

Common problems

MLOps

Role of MLflow



MLflow Tracking

Experiments, runs, parameters, and metrics

MLflow UI

The screenshot shows the MLflow UI interface. At the top, there is a navigation bar with the MLflow logo, 'Experiments', 'Models', and links to 'GitHub' and 'Docs'. Below the navigation bar, the 'Experiments' tab is selected, showing a list of experiments. One experiment, 'Default', is listed with options to edit or delete it. Another experiment, 'Dropping columns', is also listed with similar options. To the right of the experiments list, there is a search bar with the query 'metrics.rmse < 1 and params.model = "tree"'. Below the search bar, a table displays 10 matching runs. The table has columns for checkboxes, Start Time (all entries are '1 month ago'), Run Name (all entries are '-'), User ('pawel'), Source ('experiment'), Version ('505530'), Models ('-'), and MSE ('2662507.').

	Start Time	Run Name	User	Source	Version	Models	MSE
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	2662507.
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	2854716.
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	3305431.
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	3980909.
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	2686623.
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	-
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	-
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	-
<input type="checkbox"/>	1 month ago	-	pawel	experiment	505530	-	-



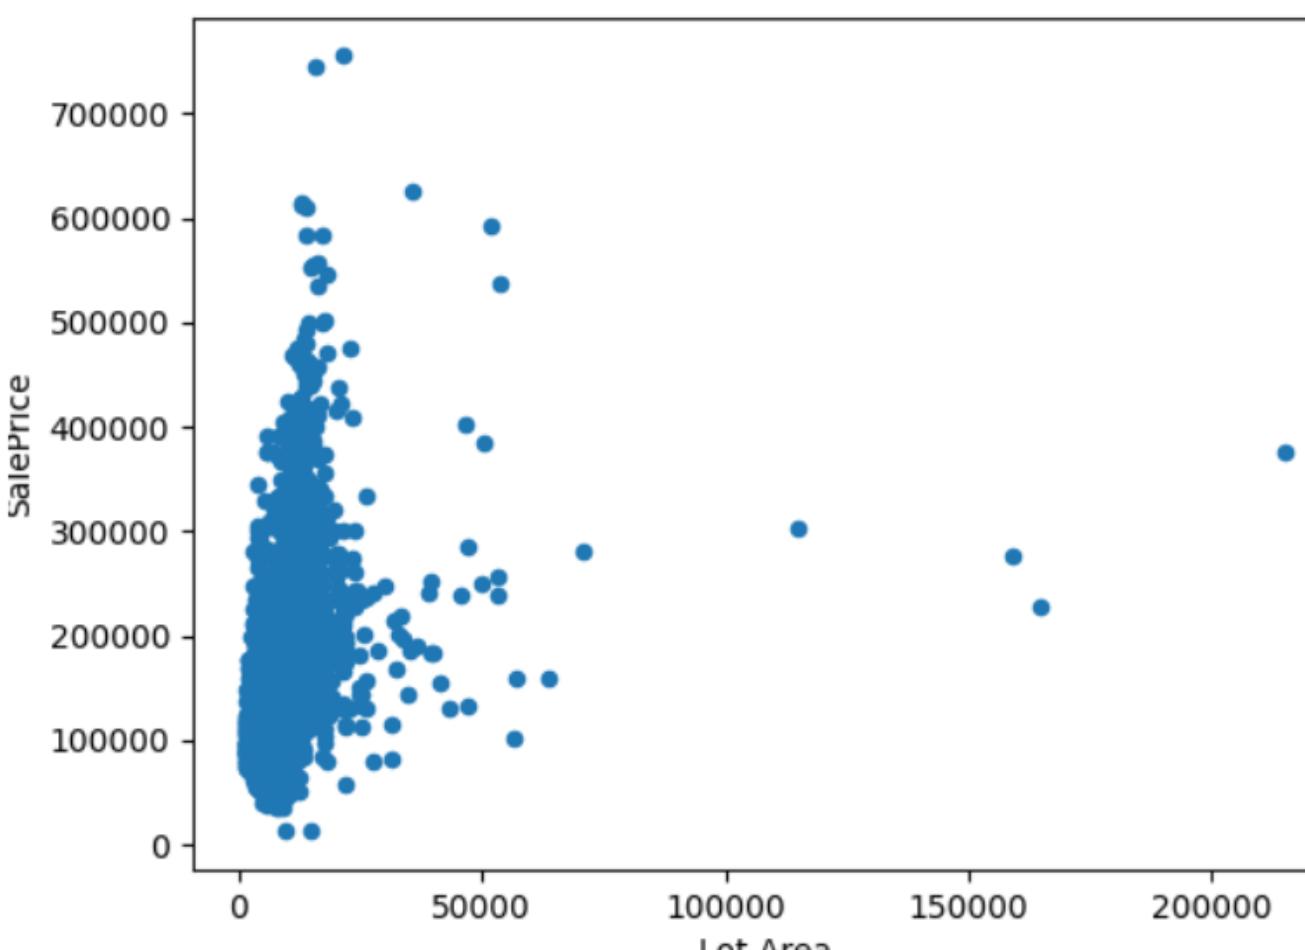
Artifacts

▼ Artifacts

- dataset.csv
- model.pkl
- plot.png

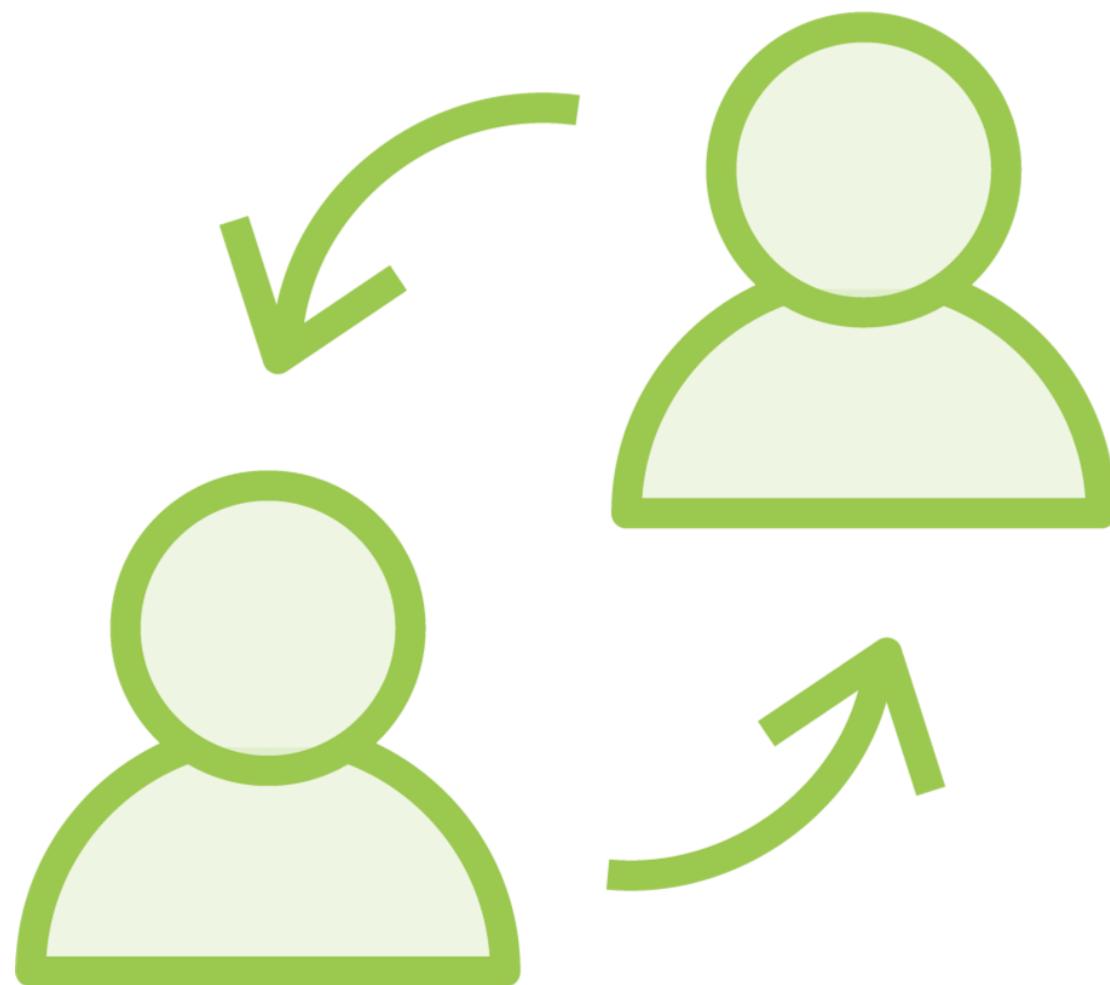
Full Path:file:///module4/mlruns/1/cf029ad774044220bdc64c96bb... ↗

Size: 24.12KB

A scatter plot showing the relationship between SalePrice (Y-axis, ranging from 0 to 700,000) and Lot Area (X-axis, ranging from 0 to 200,000). The data points show a positive correlation, with most houses having a lot area between 10,000 and 50,000 square feet and a sale price between \$100,000 and \$600,000. There are a few outliers with larger lots and higher prices.



MLflow in Teams



- MLflow architecture**
- Backend store vs. artifact store**
- MLflow server**
- Complete MLflow deployment**

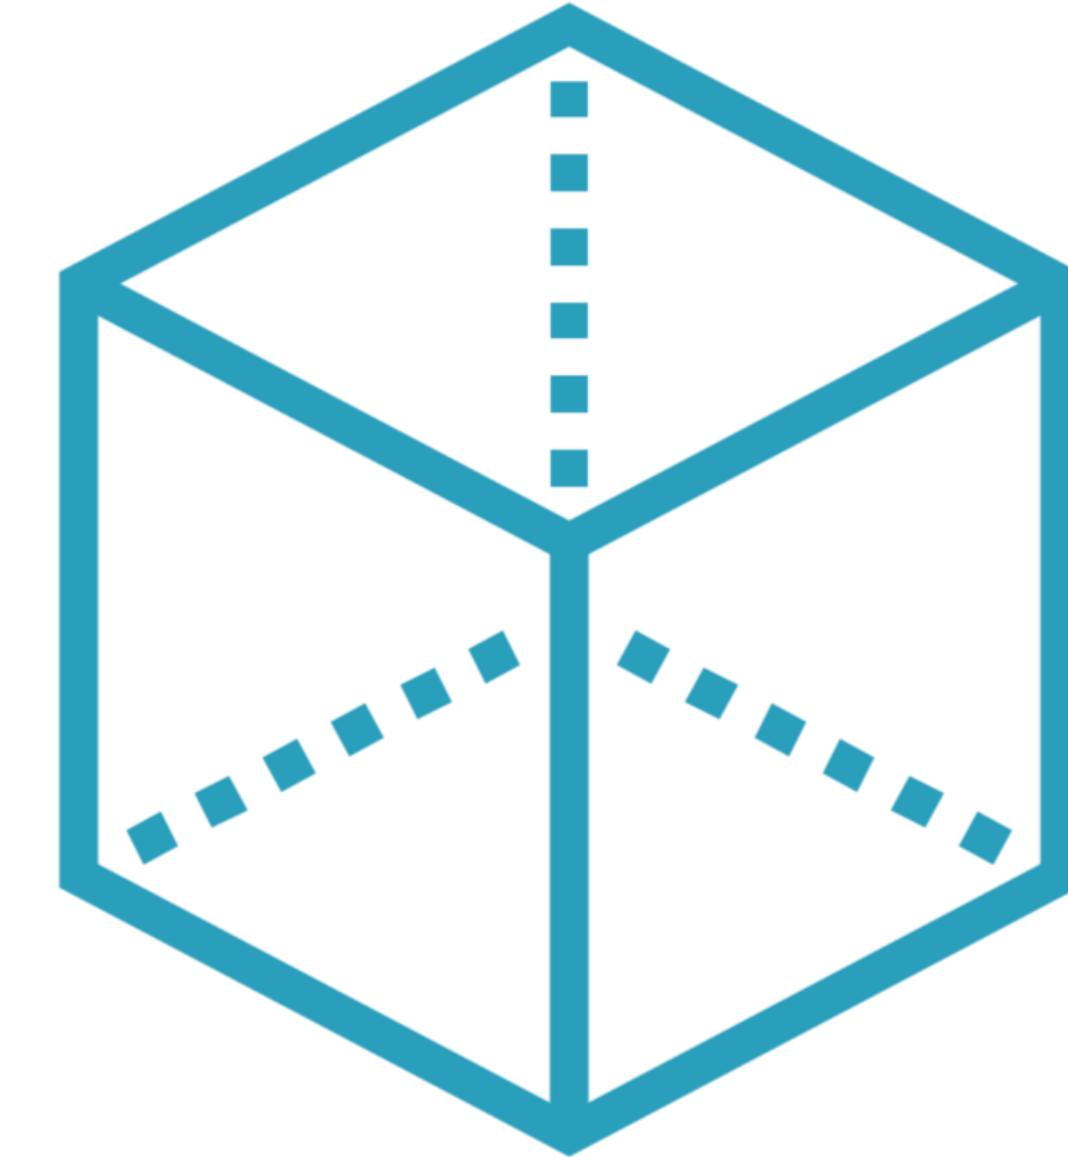


MLflow Models

Supported flavors

'Pyfunc' flavor for custom models

Running predictions in batch + real-time



Model Registry

Registered models

**Versioning and
stages**

**Running
predictions**



Thank You!

